

## Preface

"Why do you need a version control system?" ---- It's very easy to ask this question before you have learned about version control and what it can do for you. After all, RCS is just another software tool that you have to understand. These tools have a reputation for being difficult and arcane. Why give yourself more trouble than you need?

Whether or not you know it, you have probably wished you had a version control system several times in the past month. If you're at all typical, you probably have improvised a rudimentary version control system of your own. Therefore, you can understand why a version control system is necessary just by thinking about how you work. Are you nervous about what will happen when you modify a large or important file? Do you worry about what will happen if you can't reconstruct the original version of the file? Have you ever made a "safe" copy (or wished you had such a copy) before starting to edit a file?

If the answer to any of these questions is "yes," you should look seriously at what RCS has to offer. RCS really doesn't do anything more than automate what you've been doing by hand all along: it lets you reconstruct the prior versions of a file, and prevent other users from modifying your files until you are done. When you are only editing a single file, it is easy to save the old version the file before modifying it. However, this method is impractical when many files are involved or when many developers are working on the same project.

What do we mean by "automating" the version control of your project? As we've said, it quickly becomes impractical to keep "fallback" copies of old versions, track changes made, and so on. If you are involved in a large development project, you need a database to track the changes made to a file, let you reconstruct old versions easily, and record the intent and author of each modification. RCS provides such a database. It solves these management problems by efficiently storing and managing multiple versions of source files. It also gives you a way to archive the program at crucial stages, allowing you to reconstruct old versions as necessary for debugging and maintenance. In addition it provides many utilities which assist in the version management process.

Unfortunately, version control systems have an unjustified reputation for being difficult to use. This reputation has prevented them from being used more widely. If you are new to version control, this book will help to de-mystify the topic for you so that you can feel free to use RCS to solve your version

management problems. If you are already a competent user of RCS, this book will acquaint you with more advanced features that will enhance your productivity. Here are some of the book's more important features:

- We explain the underlying concepts of version control.
- We provide many examples of how to use the advanced features of RCS.
- We describe common mistakes and how to overcome them.
- We provide a detailed reference on user-supplied enhancements to RCS that are available in the public domain.

This book will therefore be particularly helpful to:

- Potential users who would like to find out what a version control system can do.
- Existing users of RCS who wish to learn how they can get more from their system.

This book can be broken down into three sections. The first section contains a chapter which introduces the basic commands and concepts of RCS. This chapter will teach you how to start controlling your files with the RCS system. It covers the fundamental features only.

The second section introduces you to some of the more advanced features of RCS itself. This section contains descriptions of identification markers, branches, symbolic names, states and general utility commands.

The third and final section describes how you can extend the functionality of RCS to make it more useful for team development efforts. Despite the fact that RCS is frequently used to control the source files for a team based project, RCS has several weaknesses when used by teams. In this section there is a chapter describing some shell scripts which you can use with RCS to make it much more suitable for use with a team based projects. We have also included a chapter describing the security problems that arise in team based projects and how them can be overcome. Last but not least there is a chapter which describes a number of RCS based systems which provide a team based version control service. All of these systems provide more functionality than RCS, but they use RCS itself for the basic version management tasks.

I would like to acknowledge the assistance of Mike Loukides who was the editor of this book. His suggestions greatly improved this book. I would also like to acknowledge the assistance of Donal Daly who originally came up with the idea for the book. Last but by no means least I would like to thank my wife Frances and my daughter Fiona for their great patience with me while I wrote the book.

## 1

## RCS Basics

"Nothing changes more constantly than the past; for the past that influences our lives does not consist of what actually happened, but of what we believe happened".

*Gerald White Johnson  
American Heroes and Hero Worship, 1943*

If you have never used RCS before, you will probably be a little puzzled by the whole concept of a version control system. This chapter we will help you by introducing the concept of version control and describing how to get started with RCS. First we introduce the files that RCS uses to maintain its databases and explain the relationship between them. We then describe how you would place a file under source management, how to retrieve a working copy of a file that is under source management, and how to retrieve an older version of your file.

## Getting Started

Version control systems like RCS store and manage different versions of text files. For each file that you manage, RCS creates two files: a *working file* and a *library file*. The *working file* is the file that you edit, compile, or even distribute. It is no different from any other text file, except that it is managed by your version control system. The *library file* is the RCS system's database about the working file. It contains the most recent version of the working file, a description of how to construct any older version, and other archival information: author, creation date, comments, and so on. You rarely need to work with the library file directly; the version control system provides tools to manage it. If you look at the library file,

you should be able to find a copy of the working file within it with little trouble. But there will be a lot of other junk. You obviously can't compile a library file, and would never send it to a customer.

The two most basic operations that any version control system performs are called *check-in* and *check-out*. The *check-in* operation stores the current contents of the working file within the library file. It also records the time of the check-in, the person performing the check-in, the file's version number, and other information. Once you have *checked-in* a version of a file, you cannot modify that version any more. You can only create new versions; you cannot change old versions. This ensures that you can accurately trace the development a program's development. You can use the *check-out* operation to retrieve and modify an old version of the file--but this creates a new version, rather than changing the older version that you've already checked-in.

In order to clarify the relationship between each library file and the corresponding working file, the system assigns it a similar name. The name of the RCS library file is generated by appending ,v to the end of the working file name. The following examples show how this works :

Table 1-1: Library File Names

<b>Working file</b>	<b>Library file</b>
<i>test.c</i>	<i>test.c,v</i>
<i>sample.h</i>	<i>sample.h,v</i>

If many files from one directory are under RCS control, the directory can become cluttered with all of the library files. In order to avoid this problem most users put the library files into a sub-directory of their own. This is often called an "archive" directory. By convention, the sub-directory used for storing RCS library files is called *RCS*. The following table show the mapping between working file names and their associated library file names when using a sub-directory.

Table 1-2: Library Files in Sub-directory

<b>Working file</b>	<b>RCS library file</b>
<i>test.c</i>	<i>RCS/test.c,v</i>
<i>sample.h</i>	<i>RCS/sample.h,v</i>

While you don't need to store files library files in a special subdirectory, we strongly recommend that you do. It will be much easier for you to work if your workspace is uncluttered. Furthermore, if you are using RCS properly, you should never have to peek at the library files.

Before you can start using RCS, you must create library files for each working file that you want to manage. Assume that you want to place a single file named *sample.c* under RCS control, and that you have already created the a sub-directory *RCS* to hold the library files.

The command for initializing a library file is **rcs -i**. Therefore you can generate a RCS library file to be associated with *sample.c* by entering the command:

```
% rcs -i sample.c
```

This automatically generates a library file named *RCS/sample.c,v*. RCS does not care whether you use the name of the working file or the name of the library file because the system can automatically convert from one form to another. Therefore you could equally well have typed the command:

```
% rcs -i RCS/sample.c,v
```

The *rcs -i* command only creates a skeleton library file. Even if the working file *sample.c* exists, its contents aren't reflected in the library file yet. In other words, you have initialized a database, but you haven't put any data in it.

Before you can go any further, you must insert an initial version of the working file into the library file. This is called *checking-in* the current contents of the working file. The RCS check-in command is **ci**. Therefore, you use:

```
% ci sample.c
```

When you check-in the first version of a file, the system issues a diagnostic message indicating that it has numbered the first version 1.1.

You frequently want to create an initial version at the same time as you create the library file, so RCS provides a convenient way to perform the two operations together. RCS will automatically create the library file if one does not already exist when you issue the *ci* command. This means that the *rcs -i* command is redundant in most situations.

When the contents of the working file are checked-in to the library file, the system deletes the working file. You don't need the working file any more; you can re-create its contents at any time by checking it out from the library file. Furthermore, if the working file still existed, it would only tempt you to make mistakes by editing it directly and circumventing source control.

To retrieve a working file, you must "check it out," as you would check a book out of a library. The *check-out* operation extracts an editable working copy from the library file. The RCS check-out command is *co*. This command re-creates the working file *sample.c* as it was immediately prior to check-in:

```
% co sample.c
```

At this point, there is no need to specify which version you want because there is currently only one version stored in the library file. Later, when you have stored several versions of the file in your library, these commands will get the most recently created version. There are several different ways to retrieve other versions, but before this can become a problem you must first create a new version and store it in the library file.

To create a new version of a file, you have to edit it. But if you try, you'll find that you can't edit the file that you've just checked out. By default, when you re-create a version using the check-out operation, the resultant working file is created in read-only mode. This reminds you that a file should not be updated without first *locking* the associated library file. RCS employs an exclusive locking scheme (only one user is allowed to lock a file at a time) to ensure that two developers can't work on the same file simultaneously.<sup>†</sup>

---

<sup>†</sup> The "locking" that RCS performs is completely independent of the file locking facilities that your operating system may (or may not) support. We only describe the locking provided by RCS.

You can lock an RCS file (and create a writable working file) by using the `-l` (lock) option with the check-out command. *e.g.*:

```
% co -l sample.c
```

After you have created a writable working file you can then use a text editor or other tool to create the new version of the working file (*e.g.*, *vi sample.c*). Once you have finished editing the file, the next step is to check-in the changed working file into the library file as a new version. The RCS check-in command is *ci*. *e.g.*:

```
% ci sample.c
```

RCS stores the current contents of the working file as a new version in the library file with an associated version number (1.2). The system then prompts you to enter a commentary: you should type in a short comment describing the changes you have made since the last check-out. The commentary can be used later by someone who is trying to understand your changes or to debug your code.

Before moving on, let's ask one more question. Why would you ever want to "check out" a file without locking it? Why would you want a read-only copy of the file? Assume that you are working with five other programmers on a large project. Each programmer is primarily responsible for one file, though he or she may need to edit other files from time to time. You obviously need to check out, locked, any file that you intend to edit. However, you just can't ignore the rest of the files: you'll need all six files for testing. And you don't want to lock these--the other developers may want need to modify them, and you don't want to prevent them from doing their work. Therefore, you want to check out all six files for the development project: your own plus the other five. But you should only "lock" the files that you plan to edit--leave the others unlocked.

## Version Numbering

Once a file is under RCS control, each version of the file has a unique version ID number. These numbers are not assigned arbitrarily; version ID numbers have a special structure and significance. You have already seen some ID numbers: the ID number 1.1 was assigned to the initial version of *sample.c* when it was checked in. After *sample.c* was checked-out, modified, and checked-in again, it was assigned the number 1.2.

Version numbers consist of positive two integers separated by a dot (*e.g.*, 3.7). The integer before the dot is called the *release number* while the integer after the dot is called the *revision number*. Therefore a version numbered 3.7 is revision 7 of release 3. The first version created is normally assigned a version id number of 1.1 (revision 1 of release 1). By default, every newly created version is assigned a version id number with the same release number as the current version but a revision number that is one higher. Thus the revision number differentiates between subsequent versions. The need for a release number is not as obvious to a novice user. We therefore introduce an example which will show you the benefit of release numbers.

The benefit of having separate release and revision numbers can not be seen by looking at a single file. Most software systems contain a number of inter-related source files which must be developed together. Release numbers only become meaningful when we look at these larger systems. In order to illustrate the benefits of the release numbering system, we will consider a sample system containing two files.

*system.c*    The file containing all routines etc.  
*system.h*    The file containing variable definitions etc.

When they were created, both files were assigned the initial version number 1.1. After a few months of development work, the company is happy with the current status of the system and ready to start selling the system to customers. At this time *system.c* has been changed 9 times and hence its current version is numbered 1.10. However, *system.h* was only changed twice in the same period and hence its current version is numbered 1.3. Development work will continue on this system after the release, creating further versions of both files. Since version 1.10 of *system.c* and version 1.3 of *system.h* are going to be compiled together and released to a large number of customers, it is important that the company should be able to remember easily that version 1.10 of *system.c* and version 1.3 of *system.h* combine to form release 1 of the system as a whole.

However, if the release level of the system were left unchanged, subsequent versions of *system.h* would be numbered 1.4, 1.5, etc., and the subsequent versions of *system.c* would be numbered 1.11, 1.12, and so on. It would be easy for the developers to forget or get confused about which versions were used to form release 1 of the system. To eliminate confusion, subsequent versions of both files should be assigned a release number of 2 to show that they belong to the development of release 2 of the system. This means that the sequence of version numbers in *system.c* will be ... 1.10, 2.1, 2.2 .... and the sequence of version numbers in *system.h* will be ... 1.3, 2.1, 2.2 ... This makes it clear that version 1.10 of *system.c* was used to build release 1 of the system because it is the highest revision number of the file with a release number 1. It is also clear that version 1.3 of *system.h* was used to build release 1 of the system as a whole because it is the highest revision number of this file with a release number 1.

While even this simple example shows the benefit of using a release based numbering convention, most software development projects contain more than two source files. As the number of files increases there is a corresponding increase in the benefit of following the release numbering convention. You should be aware that it is important to put some thought into the release/revision numbering scheme before creating a new version of any file: it is not possible to change the ID number of a version after it is created. You should also be aware that "external" release numbers (like "Version 4.25 of the Software Bugzapper") are conveniences for the public and for your marketing department, but they don't necessarily correspond to RCS version numbers. When you are developing a program, use conventions that make sense to you.

Normally RCS automatically assigns a version ID number to a newly checked-in version by incrementing the revision number of the predecessor. For example, if the previous version was numbered 1.3, the next version will be numbered 1.4. Likewise, if the current version is numbered 3.2, the next version will be numbered 3.3. You can override the default version number by specifying an alternate number for the new version with the *ci -r* command. For example the following command checks-in a new version and assigns it a version id number of 4.1.

```
% ci -r4.1 sample.c
```

If you specify a release number with the *ci -r* command (instead of a complete version id), RCS will assign the new version the lowest possible version id number from that release. For example, if the highest version id number currently stored in the library file *sample.c,v* was 4.3 the following command would cause the newly created version to be assigned the version id number 5.1.

```
% ci -r5 sample.c
```

However, if there already was a version with an id number of 5.3, the newly created version would be assigned a version id number 5.4.

You can specify any arbitrary version number when checking in a file with the *ci -r* command. The *ci -r* command is usually used to increment the release number: e.g., to go from release 2 to release 3.1. However, the only restriction that RCS enforces is that the new version number must be higher than the highest existing version number (to avoid duplicates). The following table shows how you might use the *ci -r* command:

Table 1-3: New Version Numbers Assigned

Command Used	Current Default Version Number	New Version Number
% <b>ci sample.c</b>	1.7	1.8
% <b>ci -1.9 sample.c</b>	1.7	1.9
% <b>ci -r1.3 sample.c</b>	1.7	<b>error</b>
% <b>ci -r2 sample.c</b>	1.7	2.1
% <b>ci -r2 sample.c</b>	2.3	2.4
% <b>ci -r1 sample.c</b>	2.3	<b>error</b>
% <b>ci -r3 sample.c</b>	1.7	3.1
% <b>ci -r3.3 sample.c</b>	1.7	3.3

## Retrieving Old Versions

If you perform a check-out operation on a library file with more than one version you will, by default, be given a copy of the most recently created version. The main advantage of using a version control system, however, is that you have a choice of retrieving any arbitrary version. RCS provides you with a number of ways of specifying that you want a version other than the default version.

The most straight forward way to specify that you want a different version is by giving the version id number of the version you want with the *-r* option to the check-out command. For example, the following command will retrieve version 1.1 of *sample.c* :

```
% co -r1.1 sample.c
```

When you retrieve a version using the *co* command, the system will automatically retrieve the default (most recently created) version from the relevant library file. You can retrieve another version by specifying the version id of the version required with the *-r* option flag. If you specify a release number rather than a complete version id number, the system will select the most recent version of that release number.

For example, consider a file containing six versions with the following version id numbers:

```
1.1
1.2
2.1
2.2
```



2.3  
3.1

The table below shows some of the `-r` options that you might use and which version each command would retrieve:

Table 1-4: Use of `-r` Option

<b>RCS Command</b>	<b>Version Selected</b>
<b><code>co -r2.3</code></b>	2.3
<b><code>co -r2.2</code></b>	2.2
<b><code>co -r1</code></b>	1.2
<b><code>co -r2</code></b>	2.3
<b><code>co -r3</code></b>	3.1

Without a `-r` option, `co` would of course retrieve version 3.1.

The `co -r` command is fine when you know the version ID number of the version which you need. However, you might not know exactly which version of the file you are looking for. Luckily RCS supports several mechanisms for specifying which version you require.

For example, you might know that you want to retrieve the version which was current on a particular date. The `co -d` command lets you specify that you are only interested in versions created before a certain date. For example, the following command checks-out the version of *sample.c* that was created closest to, but before, the 1st of September 1990:

```
% co "-d1-sept-90" sample.c
```

The date specified with the `-d` option can be specified in *free format* as described in appendix D. You must be careful to enclose the date string in quotation characters if it contains any spaces.

You can also select the most recently created version that was created by a particular user by using the `co -w` command. For example, the following command retrieves the latest version of the file *sample.c* that was created by the user *fred*:

```
% co -wfred sample.c
```

The various version selection options can be used in conjunction with each other. For example, the following command will retrieve the version of *sample.c* that was created closest to (but before) 1st of September 1990, that was authored by *fred*:

```
% co "-d1-sept-90" -wfred sample.c
```

You can also select versions based upon their state, symbolic name and/or branch number. However, we will postpone discussion of these features until later in the book.

## Locking

While you may use RCS to manage multiple versions of files that you are developing yourself, RCS is more commonly used by teams of programmers to store files that are being jointly developed by the team members. RCS uses locking to ensure that work being done by one team member is not in conflict with work being done by another team member. You must lock a file before starting to update it. A lock ensures that no one else is working on the same file at the same time.

When a number of people are working together on a project they often need to jointly develop some of the source files. This leads to the possibility that two (or more) of the developers might make conflicting updates to the library copy of the file. There are two distinct ways in which to deal with conflicting updates; we can either attempt to resolve the problem after it happens or we can try to avoid having conflicting updates happen in the first place. RCS tries to ensure that conflicting updates never take place by implementing an exclusive locking scheme.

The problem of conflicting updates can easiest be understood by looking at an example of how a conflicting update might arise. Suppose there were two programmers Jim and Mary both working on different sections of a project but jointly developing the user manual. On monday Jim checks-out the most recent version (version *1.1*) of the manual source file in order to insert a description of some routines he has just coded. These descriptions are rather lengthy and hence he does not have the modified version of the user manual ready until thursday.

On tuesday, while Jim is still working on the descriptions of his new routines, Mary notices some minor spelling changes in the manual and decides to correct them. She checks-out the most recent version (still version *1.1*) of the user manual to make the corrections. She makes these corrections and checks-in her corrected copy of the manual. The system assigns it version number *1.2*.

When Jim checks-in his copy of the user manual containing the description of his routines, it is assigned version number *1.3*. Since Jim started working with version *1.1* of the document the spelling errors from version *1.1* are in version *1.3* the document. Hence the effect of Mary's spelling corrections (the update from *1.1* to *1.2*) have been effectively lost in version *1.3*.

This problem is sometimes called the lost update problem because Mary's update has been lost. This problem is also called the conflicting update problem because the two updates are in conflict with each other.

When dealing with conflicting updates, as in many other aspects of life, avoidance is better than a cure. The key to avoiding conflicting updates is to ensure that only one user is updating a file at any given time. In RCS this is accomplished by locking. RCS insists that you must lock a file before updating it. Since RCS ensures that only one user may lock a file at any given time, this means that only one user can be updating the file.<sup>†</sup>

When a user checks-out a copy of a file from the version control system the syntax of the command indicates whether they want either a read only copy or a write/read copy.

If the user intends to create a new revision, they will request a write/read copy. The system will attempt to lock the file on behalf of the user before retrieving a write/read copy. If the file is already locked by another, user the check-out request will fail because it is not permissible for two users to simultaneously

<sup>†</sup> In a later chapter we will describe branches; this is a feature which allows you to have two or more separate lines of development within one file. When you are using development branches, RCS allows that two different users may have locked different branches of the same file. This is safe because updates to different branches will not be conflicting.

lock the same file.

In the earlier example, Jim would have locked the source file for the user manual when he checked-out version one. This would mean that Mary's request to check-out a write/read copy on tuesday would have been denied thus forcing the two updates to be carried out in series rather than in parallel. Jim's lock would have been automatically relinquished on thursday when he checked-in the new version of the file, Mary would then have been free to proceed with her changes.

A user may check-out a read only copy of a file regardless of whether the file is locked or not. In order to remind users that they should not be editing a copy of a file that was checked-out without a lock, the UNIX write permissions are normally turned off even for the owner of the file if it is checked out without a lock.

Even with the locking mechanism it is always possible that people will work around it and create a conflicting update. In this case we are left with the problem of how to resolve the conflicting update. This is done by merging the lost update into the most recent version.

If we take our earlier example, we would like to create a new version (version 1.4) which would contain both the description of new routines (from version 1.3) and the spelling corrections (from version 1.2). Special merge utilities are provided to automatically accomplish this task. So long as the update from version 1.1 to 1.2 affected different parts of the file from the parts affected by the update from version 1.1 to 1.3 it is possible to automatically make a fairly intelligent merge of the two updates by using the *rcsmerge* command which will be described later in this book. This command merges the effect of both changes by including the changed copy of each section. However, if the two updates affect the same section of the file, it is not possible to guess which update should have preference and the conflict between these two updates must be resolved manually.

Normally a lock is only acquired as part of a check-out operation and it is relinquished as part of a check-in operation. However, this might prove to be too restrictive for many situations. As a result RCS provides a mechanism for relinquishing a lock without creating a new revision, acquiring a lock without getting a copy of the file and even releasing the lock belonging to another user.

The *rcs -l* command will lock a file without checking-out a working file. This is useful if you check-out a read-only version of the file, and sometime later decide that you need to edit it. Rather than performing another check-out operation, you can use *rcs -l* command shown below to "lock" the copy of the file you already have.

```
% rcs -l sample.c
```

Normally the *rcs -l* command will lock the default version of the RCS file (*i.e.* the version that would normally be retrieved by a *co* command). You can also specify that you want to reserve any other version by giving a version number with the *rcs -l* command. (It is necessary to lock a version other than the default version when you want to create a branch as will be described in a later chapter.) For example, the following command locks version 1.3 of *sample.c*.

```
% rcs -l1.3 sample.c
```

Likewise, you can unlock a file without checking-in a working file using the *rcs -u* command. This could be useful if you started to make a change and then realized that the change was not necessary; you do not want to check-in the file with the partial update. You can delete your changed copy of the file, but you must also relinquish the lock which you were granted when you executed the *co -l* command. The following command will unlock whichever version of the file *sample.c* you have locked.

```
% rcs -u sample.c
```

It is possible that you could have locked more than one version of a single RCS file; this would happen if you were working on two different branches of the file simultaneously. In this case you will need to tell RCS which of the locks you want to relinquish with the *rcs -u* command. For example, the following command unlocks version 1.3 of *sample.c*.

```
% rcs -u1.3 sample.c
```

If you wish to update a file that is already locked by somebody else, you must find out which user "holds the lock": i.e., find out which user has locked the file. Normally you would then contact that user and ask him or her to relinquish the lock. However, if you can't contact the other user, RCS will let you "break" the lock. This may be useful if a co-worker leaves for vacation and forgets to check-in the files that he or she is working on.

Since the primary purpose for locking is to ensure that no two users are unknowingly working on the same file, RCS adopts the policy that it is permissible to break someone else's lock, provided the original locker is given a clear indication that this has happened. If you break a lock belonging to another user, the system will ask you why you are breaking the lock. RCS then sends a mail message, including your reply, to the user holding the lock, informing him or her that the lock has been broken.

The following example shows how this works:

```
% rcs -u1.3 sample.c
RCS file: sample.c,v
Revision 1.3 is locked by fred
Do you want to break the lock? [ny](n): y
State the reason for breaking the lock:
(terminate with ^D or single '.')
>> I cannot wait for you to come back
>> from vacation
>> .
1.3 unlocked
done
%
```

RCS then sends the following mail message to *fred* to let him know that you have just broken his lock:

```
Date: Thu, 13 Sep 90 14:57:35 +0100
From: mary
To: fred
Subject: Broken lock on sample.c,v

Your lock on revision 1.3 of file sample.c,v
has been broken by mary for the following reason:

I cannot wait for you to come back
from vacation
```

In our example, we explicitly specified that we wanted to unlock version 1.3 of *sample.c,v* by specifying the version number with the command (*rcs -u1.3*). Some releases of RCS break the lock on the most recent version of the file if you do not give the version number explicitly. Other releases of RCS refuse to break the lock if you do not give an explicit version number.

Whenever you check-in a working file, the system automatically deletes the working file and deletes the lock you held on the file. You often want to continue working on the next version of a file immediately after checking-in the current version. In this case, you can use the `ci -l` command to check-in the file without relinquishing your lock or deleting the working file. For example, the command:

```
% ci -l sample.c
```

is equivalent to a check-in followed immediately by a check-out. The following two commands have the same effect:

```
% ci sample.c
% co -l sample.c
```

At other times, you might want to maintain a read-only copy of the working file (e.g. for compilation) but you do not want to maintain your lock because you are not going to update the file further. In this case, you can use `ci -u` to check-in the file and relinquish your lock without deleting the working file. For example, the command:

```
% ci -u sample.c
```

is equivalent to the following two commands.

```
% ci sample.c
% co sample.c
```

Locking is a valuable tool for avoiding conflicting updates when many users are jointly working on a file. However, it can sometimes be a pain. If you are the only person updating a file, you may feel that locking is nuisance you can live without. RCS gives you the option of turning off the strict implementation of the locking feature. When locking is strictly enforced, you must hold a lock on a file before you can check-in a new version. If strict locking is not enforced, you don't need to hold a lock on files that you already own. However, even when strict locking is not being enforced, you may not check-in a new version of a file which is locked by someone else.

Strict locking is normally the default. However, your system administrator can change this default when he or she first compiles RCS. If non-strict locking is desired, the macro `STRICT_LOCKING` should be defined to be equal to 0 in the source file `rcsbase.h`. If strict locking is desired, the macro `STRICT_LOCKING` should be defined to be equal to 1.

You can also enable or disable strict locking on a file-by-file basis by using the `rcs -L` or `rcs -U` commands respectively. For example, the following command enables strict locking for `sample.c,v`.

```
% rcs -L sample.c,v
```

Likewise, the following command disables strict locking for `sample.c,v` and `sample.h,v`.

```
% rcs -U sample.c,v
```

## NOTE

Because of the way that RCS handles updates to library files, the owner of an RCS file is always the user who made the last update, rather than the user who created the file. Hence, it is the person who last updated the file who may be exempt from needing to hold a lock, not the person who created the file.

## Traceability Information

You will often find bugs in one version that did not exist in a previous version. The obvious question to ask is: "what has changed between these two versions?" The *rcsdiff* command (which is described later in the book) can be used to find these changes. However, before you can fix the bug, you should be able to answer the second question: "what did the author intend when making these changes?" You can't just assume that the changes were mistakes. The author was no doubt trying to do something intelligent, and if you don't understand the intent, you will only introduce new bugs. There is no way to answer the second question automatically. Instead, the system stores some traceability information to assist you. RCS stores a log message with each version. It also stores a general description of the file when you first put it under control.

Whenever you check-in a new version of a RCS file, the system prompts you to enter a log message describing the change you have just made. This log message is stored in the RCS file along with the new version itself. You should enter a clear message describing what changes you made and why you made these changes: the log messages can serve as a valuable reminder to someone who is later trying to understand what you have done. Don't just enter a message like "fixed a bug"--give a description that will be useful to someone who needs to work with the program after you've moved to another project. The log message may span several lines. You indicate the end of the log message by typing a line that contains a single period (as the first character) or a line containing a single CTRL-D. The following example shows a typical *ci* operation:

```
% ci sample.c
RCS/sample.c,v <-- sample.c
new revision: 1.3; previous revision: 1.2
enter log message:
(terminate with ^D or single '.')
>> added an input checking routine to eliminate the
>> input consistency bug reported by sally
>> .
done
%
```

If you check-in a number of files with one *ci* command, you may want to give the same log message to each of the newly created versions. To save you from typing the same log message repeatedly, the *ci* command gives you the option of reusing the first file's log message for all subsequent files. The following example shows how this might be used:

```
% ci sample.c routines.c
RCS/sample.c,v <-- sample.c
new revision: 1.4; previous revision: 1.3
enter log message:
(terminate with ^D or single '.')
>> changed these files so that they now use double
>> precision math functions
>> ctrl-D
done
RCS/routines.c,v <-- routines.c
new revision: 1.3; previous revision: 1.2
reuse log message of previous file? [yn](y): y
done
```

You may prefer to give the log message on the command line rather than using the dialogue format shown above. You can do this by using the *ci -m* command. For example, the following command will check in a new version of the file *sample.c* with the log message "fixed fp convergence problem".

```
% ci "-mixed fp convergence problem" sample.c
RCS/sample.c.c,v <-- sample.c.c
new revision: 1.5; previous revision: 1.4
done
%
```

If the log message contains spaces, you should be careful to put quotation marks around the `-m` option, as in the example. If we had omitted the quotation marks, the shell would have interpreted the words "fp" and "convergence" as file names, rather than as a continuation of the log message.

The `ci -m` command is normally used within shell scripts. Because of its syntax, it is really only appropriate for a short (one-line) comment. However, you can give a multi-line log message with the `ci -m` command by putting the ``` character at the end of each line. This tells the shell that the command has not ended yet. For example:

```
% ci "-mthis log message\
spans two lines" sample.c
```

will check-in a new version of the file *sample.c* with the log message:

```
this log message
spans two lines
```

Whenever you create a new RCS file, either with the `rcs -i` command or with the `ci` command, the system prompts you to enter a descriptive text. This descriptive text describes the entire RCS file and should not be confused with the log message. The log message is given whenever you check-in a new version, and only describes a particular version of the file; the first version of a RCS file is always given the log message "Initial revision." The description is only given once, and describes the purpose of the file as a whole.

To enter the descriptive message, follow the same rules that you use for the log message. Type as many lines as you like, terminated by a line containing a single period or CTRL-D, as shown in the following example:

```
% ci f_handle.c
RCS/f_handle.c,v <-- f_handle.c
initial revision: 1.1
enter description, terminated with ^D or '.':
NOTE: This is NOT the log message!
>> This file contains the code to implement
>> the file handling routines
>> .
done
%
```

Normally, RCS only asks you to enter a descriptive message when it first creates the library file. However, you don't always know exactly what a file will contain when you first create it. Over time, your initial description may become outdated or incorrect. You can force the system to prompt you for a new descriptive text (and discard the old one) by using either the `rcs -t` or `ci -t` commands. For example, both of the following two commands discard the original description and create a new one.

```
% rcs -t f_handle.c
RCS file: RCS/f_handle.c,v
enter description, terminated with ^D or '.':
NOTE: This is NOT the log message!
```

```
>> This file contains the enhanced file handling routines.
>> .
done
```

or

```
% ci -t f_handle.c
RCS/f_handle.c,v <-- f_handle.c
new revision: 1.2; previous revision: 1.1
enter description, terminated with ^D or '.':
NOTE: This is NOT the log message!
>> This file contains the enhanced
>> file handling routines.
>> .
enter log message:
(terminate with ^D or single '.')
>> Added enhanced routines and updated description
>> .
done
%
```

If you want to include a lengthy description of the file, you might prefer to prepare the description with your favorite text editor. Use the editor of your choice to create a text file containing your description, and then use the `-t` option to `rcs` or `ci` to specify your description file. If you name a description file with the `-t` option, RCS uses the description from the file. For example, the following command copies the descriptive text contained in the file *txtfile* into the RCS file *RCS/f\_handle.c,v*.

```
% rcs -txtfile f_handle.c
```

Similarly, the next command copies the descriptive text contained in the file *txtfile* into the RCS file *RCS/f\_handle.c,v*; but it also checks-in a new version at the same time.

```
% ci -txtfile f_handle.c
```

## Specifying File Names

All the RCS commands allow you to specify either the name of the working file or the name of the corresponding library file. However, most RCS commands need to know the name of both files. It is simple for the system to recognize which of the files you have specified because RCS file names end in `,v` while working file names do not.<sup>†</sup> And given that it has one file name, it is easy for RCS to compute the other. RCS has a few simple rules for generating library file names from working file names and vice-versa. We'll start by explaining these rules. For simplicity, all of our examples use the `co` command to illustrate the relationship between working files and library files; the other RCS commands all use the same rules.

If the working file is to be created in the current directory and the RCS file is stored in the RCS sub-directory, the conversion is simple. If you issue the command:

<sup>†</sup> Because of the file naming convention used, RCS cannot be used to manage files whose names end in `,v`. However, this is not a very common file ending.



```
% co sample.c
```

RCS recognizes that you have given the name of the working file and not the library file. Hence it first looks for a library file named *RCS/sample.c,v*. If this file does not exist, it then looks for a library file named *sample.c,v*. Likewise, if you issued the command:

```
% co RCS/sample.c,v
```

RCS recognizes that you have specified the name of the library file. Hence it calculates that the name of the working file should be *sample.c*. In order to prevent confusion and give you a record of what it actually did, the *co* command prints an informative message like this:

```
RCS/sample.c,v --> sample.c
```

This means that the working file *sample.c* was checked-out from the library file *RCS/sample.c,v*.

If you are working as part of a development team, you probably keep your copy of the working file in your local directory, but store the library file in a public directory that is accessible by all team members. Therefore, if you give a full directory specification for the library file, RCS still assumes that it can find the working file in your current working directory. For example, the following command checks-out a version from the library file */usr/src/RCS/sample.c,v* into the working file *sample.c*, which is located in the current directory:

```
% co /usr/src/RCS/sample.c,v
```

However: If you specify a working file name in another directory, RCS looks for the library file in the same directory as the working file, rather than in the current directory. For example, the following command checks out a version from a library file named either */usr/src/RCS/sample.c,v* or */usr/src/sample.c,v* into a working file named */usr/src/sample.c*:

```
% co /usr/src/sample.c
```

Now, think about what happens when the library file and the working file are in different directories, and neither is the current directory. In this case, RCS can't figure out the relationship between the two filenames. You have to specify the full path names of both the working file and the library file explicitly. For example, the following command checks-out a copy of the library file */usr/src/RCS/sample.c,v* into the working file named */usr/alt/sample.c*:

```
% co /usr/src/RCS/sample.c,v /usr/alt/sample.c
```

While the working file and the library file may be in unrelated directories, the working file and the library file must have similar file names. For example, if the working file is named *sample.c* the library file must be named *sample.c,v*. You can't have a library file named *lib-sample.c,v*, or any other permutation.

The *co* command (and any other RCS command) can operate on several files at the same time. If you give two different file names with the same *co* command, RCS assumes that you want to perform two separate check-out operations. For example, the following command checks-out a version of the library file *RCS/sample.c,v* into a working file named *sample.c*. In addition, it checks-out a copy of the library file named *RCS/diff.c,v* into the working file named *diff.c*:

```
% co sample.c diff.c
RCS/sample.c,v --> sample.c
RCS/diff.c,v --> diff.c
```

So far, we have been assuming that you will always use *co* to check out a version into a working file. However, there are many situations in which you don't really want a copy of the file--you only want to look at the file on the screen. For example, you may want to browse through some code in another version without disturbing the working copy that you already have. If you would prefer to read a selected version of your RCS file on your terminal rather than checking out an actual working file, use the *co -p* command. The following command displays the default version of the RCS file *RCS/sample.c,v*, but does not affect the current contents of the working file *sample.c*:

```
% co -p sample.c
```

You can use this option in conjunction with I/O re-direction to store the selected version into some file other than the working file. For example, the following commands puts a copy of the latest version of release two of *RCS/sample.c,v* into the file *release2*, and a copy of the latest version of release three of into the file *release3*. The contents of the working file *sample.c* will not be affected:

```
% co -p -r2 sample.c > release2
% co -p -r3 sample.c > release3
```

You can also "pipe" the output of *co -p* into a print spooler to get a hard-copy.

## Generic Options

All RCS commands issue diagnostic messages that describe exactly what they are doing. These diagnostic messages are normally useful, but there are many situations in which they are an annoyance. For example, if you are writing a shell script that uses RCS commands, you may not want the RCS diagnostic messages to be cluttering up the screen. Invoking any RCS command with the *-q* option (which stands for "quiet") suppresses its diagnostic messages. The following example shows the difference between *co* and *co -q*:

```
% co sample.c
RCS/sample.c,v --> sample.c
revision 1.5
done
% co -q sample.c
%
```

Both commands have the same effect. The only difference between the two is that the first command explained what it was doing, while the second gave no explanatory messages. Note that the *-q* option flag only suppresses diagnostic messages. Error messages cannot be suppressed. If something goes wrong (for example: if *co* can't find the file you want to check out), you will get an error message from *co*, whether or not *-q* is present.

Both the *co* and *ci* commands have safeguards against accidental misuse. The *co* command always asks for confirmation before overwriting a writable working file. A writable working file might contain modifications which have not been checked-in to the RCS file. *Co* does *not* check before overwriting a read-only file. Because a read-only file cannot be modified, *co* considers it safe for deletion.

The *ci* command always asks for confirmation before checking-in a working file which has not been modified. There is little point in using RCS to store a series of identical versions to a file. Granted, there is no great harm in storing a series of identical versions, either--it wastes a little disk space and makes *rlog* output more unwieldy.

When either the *co* command or the *ci* command asks you for confirmation, the default answer to the question is "no"; this means that if you type a carriage return, the system will abort the command. The following two commands illustrate how this works:

```
% co -l sample.c
RCS/sample.c,v --> sample.c
revision 1.5 (locked)
writable sample.c exists; overwrite? [ny](n): y
done
% co -l sample.c
RCS1/sample.c,v --> sample.c
revision 1.5 (locked)
writable sample.c exists; overwrite? [ny](n):
co warning: checkout aborted.
% co sample.c
RCS1/sample.c,v --> sample.c
revision 1.5 (locked)
writable sample.c exists; overwrite? [ny](n): n
co warning: checkout aborted.
%
```

Here's a similar sequence with *ci*:

```
% ci -l sample.c
RCS/sample.c,v <-- sample.c
new revision: 1.6; previous revision: 1.5
File sample.c is unchanged with respect to revision 1.5
checkin anyway? [ny](n): y
enter log message:
(terminate with ^D or single '.')
>> Made no changes to the working file
>> .
done
% ci sample.c
RCS/sample.c,v <-- sample.c
new revision: 1.7; previous revision: 1.6
File sample.c is unchanged with respect to revision 1.6
checkin anyway? [ny](n): n
checkin aborted; sample.c deleted.
%
```

If you are certain that you are issuing the correct command, you can use the *-f* option with any RCS command to force it to proceed without asking for confirmation. The following two commands illustrate the effect of the *-f* option: you will notice that the second *co* command proceeds without asking for confirmation because the *-f* option was used.

```
% co -l sample.c
RCS/sample.c,v --> sample.c
revision 1.6 (locked)
writable sample.c exists; overwrite? [ny](n): n
co warning: checkout aborted.
% co -l -f sample.c
RCS/sample.c,v --> sample.c
revision 1.6 (locked)
done
%
```

The *-f* and *-q* options are somewhat inter-related. Whenever RCS asks for confirmation, the default answer is "no". If you use the *-q* option, RCS won't ask whether it should continue or not; it will always assume that you do not want to continue in questionable circumstances. In other words, *-q* implies the

opposite of *-f*.

The following sequence of commands show how the *-f* and *-q* options are inter-related. You will notice that the second *co* command aborts without asking you, because the *-q* option implies that a "no" answer is assumed to all questions. However, the third *co* command overwrites the working file without asking you, because the *-f* option implies that a "yes" answer is assumed to all questions.

```
% co -l sample.c
RCS/sample.c,v --> sample.c
revision 1.6 (locked)
done
% co -l -q sample.c
co error: writable sample.c exists; checkout aborted.
% co -l -f -q sample.c
% ci -q sample.c
ci warning: checkin aborted since sample.c was not changed; sample.c deleted.
%
```

## 2

## Identification Markers

"The sense of identity provides the ability to experience one's self as something that has continuity and sameness, and to act accordingly."

*Erik Homburger Erikson  
Childhood and Society, 1950*

RCS helps you to track different versions of source files. However, the ability to track source files would be of little use if you, the programmer, didn't also have access to the information stored in the RCS database. You often need to know the version number of the file you're looking at; the author; when it was last checked-in; and so on. Therefore, RCS provides "identification markers" which can semi-automatically place identification information into the working copy of your files.

For example, if you want the version id number to appear in the source file, you put the special identification marker "\$Revision\$," into the initial version of the file which is to be controlled by RCS (the working file). Identification markers are usually placed within comment strings. For example:

```
/* this is version $Revision$ */
```

Now, let's assume that you check-in the initial version (1.1) of the file with this identification marker. Later, when you check-out the working file, the system automatically replaces the above line with

```
/* this is version $Revision:1.1$ */
```

If you later create version 1.2 of the file but leave this line unchanged, the system will automatically update the identification marker to:

```
/* this is version $Revision:1.2$ */
```

It is often a good idea to include identification markers in executable files as well as in the source file. The exact way in which you do this depends upon the programming language you use. In the C programming language you can ensure that an identification text is inserted into the executable file by including the identification marker in a statically declared string. Here's how:

```
static char RCS_ID[] = "version $Revision$ of file.c";
```

When you compile this file, the identification string will be placed in an initialized but otherwise unreferenced character array within the object module. Note that this string must be **static**. When you link the program, identification strings from all of its components will be contained in the executable.

But how do you look for identification markers within an object file or an executable? You can't just scan an executable as you can scan a source code file. One alternative would be to use the *strings* command, but this will usually give you a lot of garbage instead of the markers that you're interested in. RCS provides something better. It provides a special utility to locate and print for RCS identification markers in files. This utility is called *ident*.

Some executables are compiled from many different source files; by placing the appropriate identification markers in each source you can identify which version of each source file was used to generate the executable file. The following example shows how this can be done for the system containing the two files *system.h* and *system.c*. The identification lines placed in the original source files would be

Table 2-1: Original Identification Lines

File	Identification line
<i>system.c</i>	<code>char c_id[]="\$RCSfile\$ \$Revision\$";</code>
<i>system.h</i>	<code>char h_id[]="\$RCSfile\$ \$Revision\$";</code>

After these files are checked-in, the identification lines would be changed to:

Table 2-2: Expanded Identification Markers

File	Identification line
<i>system.c</i>	<code>char c_id[]="\$RCSfile: RCS/system.c,v \$ \$Revision: 1.10 \$";</code>
<i>system.h</i>	<code>char h_id[]="\$RCSfile: RCS/system.h,v \$ \$Revision: 1.3 \$";</code>

When you perform a compilation on these files, the relevant strings will be included in the resultant binary file *system*. You can identify which versions of the source file have been used by giving the command:

```
% ident system
$RCSfile: RCS/sample.c,v $ $Revision: 1.10 $
$RCSfile: RCS/sample.h,v $ $Revision: 1.3 $
```

This clearly indicates to you that you have used version 1.3 of *sample.h* and version 1.10 of *sample.c*. For a small system made up of two files, you might be able to keep this information in your head. But if you are working on a large system with many files, you probably won't be able to remember which version of each file you are using. Identification markers are even more crucial if you are developing an application that will be shipped to some external users. Assume that a customer reports a bug. You can find out immediately which files were used to build his copy of the program by asking him to run *ident* on his system. If the customer doesn't have these utilities available, you can still ask him to put his executable on tape and ship it back. When you are debugging, there is nothing more important than knowing *exactly* what files you should be working with.

## The Complete List

Up to this point we have only described the \$Revision\$ identification marker. However, there is a lot more information about a version that you might like to store in the file other than just its revision number. Hence, RCS provides you with a large number of identification marker strings. The following table lists all available RCS markers:

Table 2-3: RCS identification Markers

Marker	Expands to
\$Author\$	The username of the person who created (checked-in) this version.
\$Date\$	The date and time that this version was created (checked-in).
\$Header\$	A complete identifier containing the full path name of the RCS file, the version id number, the creation date of this version, the author, the state and the person who has locked this file (if it is locked).
\$Id\$	Same as the <code>\$Header\$</code> marker, except that the RCS file name is given without the full path name. Not supported by all RCS versions.
\$Locker\$	If this version is locked, it expands to the username of the person who holds the lock. If this version is not locked, this marker is "empty."
\$Log\$	This is a special marker which expands to the RCS file name followed by a list containing the Revision number, the author, the creation date and the log message for each version that was checked-in. This marker behaves differently than all of the others. We will describe it in more detail below.
\$RCSfile\$	The name of the RCS file (e.g., <i>sample.c,v</i> ).
\$Revision\$	The version id number of this version.
\$Source\$	The full path name of the RCS file e.g. <i>/users/fred/RCS/sample.c,v</i> .
\$State\$	The state of this version.

The first letter of each marker is uppercase; the other letters are almost always lowercase. The `$RCSfile$` marker is the only exception to this rule. It may be obvious that you need to spell the markers correctly, but it's not so obvious that they must be capitalized correctly. Like most other UNIX tools, RCS is case-sensitive, and will ignore any markers without correct capitalization.

## The Change Log

The `$Log$` marker behaves differently from other markers in a number of ways. RCS replaces the `$Log$` marker with all of the log messages that have been given for this file. Over time, the `$Log$` marker grows into a complete history of the file's development. In this way, it is fundamentally unlike the other markers, which only reflect information about the most current revision. For example, the initial version of the working file *sample.sh* would contain:



```
# $Log$
```

After being checked-in to RCS, this might be updated to:

```
# $Log:      sample.sh,v $
# Revision 1.1  90/09/18  11:25:50  fred
# Initial revision
#
```

After a new version is created this might in turn be updated to :

```
# $Log:      sample.sh,v $
# Revision 1.2  90/09/18  11:26:11  fred
# A long commentary describing the changes
# which spans multiple lines
#
# Revision 1.1  90/09/18  11:25:50  fred
# Initial revision
#
```

Log information about version 1.1 is included in version 1.2 of the file, in addition to the information about 1.2 itself.

The `$Log$` marker is also the only marker that spans multiple lines. All other identification markers fit onto a single line. Because working files normally contains source code, you will usually put your identification markers inside of comments to ensure that they do not affect the meaning of the code. However, this strategy won't work for `$Log$` because many programming languages do not allow multi-line comments. For example, C allows multi-line comment, but shell scripts and FORTRAN don't. How can you use `$Log$` within a shell script without making it incorrect?

RCS tries to handle this problem for you. When it expands the `$Log$` marker, it precedes each line of the expanded text with a special "comment identification string." RCS guesses which comment identification string is appropriate by looking at the suffix of the working file's name (the part of the file name after the dot). The name of the working file in our example is *sample.sh*, the suffix of this file name is *.sh*; hence RCS guessed that it contained a shell script and used the string " # " as the comment identifier. The following table shows the various suffixes that RCS will recognize and the default comment identifier string it will use for each suffix.

Table 2-4: Comment Characters used in the Change Log

Suffix	Language Assumed	Comment Identifier
.c	C	*
.csh	c-shell	#
.e	efl	#
.f	FORTRAN	c
.h	C-header	*
.l	lex	*
.mac	DEC-macro	;
.me	n/troff -me macros	\"
.ml	mocklisp	;
.ms	n/troff -ms macros	\"
.p	pascal	*
.r	ratfor	#
.red	psl/rlist	%
.s	assembler	#
.sh	Bourne-shell	#
.sl	psl	%
.y	yacc	*
.ye	yacc-efl	*
.yr	yacc-ratfor	*
no suffix	unknown	#
unknown suffix	unknown	none

If you do not like the default comment leader selected by the system, you can override it by using the *rcs* *-c* command. For example, the command

```
% rcs -c"comm" sample.sh
```

causes the string "comm" to be inserted at the beginning of all future change logs in the file *sample.sh*. If we now create a new version the change log will be updated to

```
# $Log:      sample.sh,v $
comm Revision 1.3  90/09/18  11:45:26  fred
comm commentary supplied with new version
comm
# Revision 1.2  90/09/18  11:26:11  fred
# A long commentary describing the changes
# which spans multiple lines
#
# Revision 1.1  90/09/18  11:25:50  fred
# Initial revision
#
```

The changed comment leader is only used for the versions that are created *after* you give the *rcs -c* command, not for versions that have already been checked-in. When the working file is being checked-out the system merely updates the part of the change log relating to the current version; it leaves the rest of the change log untouched. (As a consequence, if you want to delete certain versions from the change log, you can simply remove the appropriate lines from the working file. Once removed, they won't reappear.)

The *rcsit* command

Many software companies put a standard format header at the top of all their source files. This header typically contains such information as the file name, version number, creation date, change history *etc.* . All of this information can be automatically generated by using RCS identification markers. However, you might consider it a laborious and mechanical task to insert the correct identification markers at the start of all the source files. Luckily, Michael Cooper from the University of Southern California has gone to the trouble of developing a special utility named *rcsit* which you can use to automatically insert the correct header information at the beginning of your files.

The following example shows how you would use *rcsit*. Initially you have a file *sample.c* with the following contents.

```
main()
{
    printf("hello world0");
}
```

You issue the command:

```
% rcsit sample.c
```

the file is updated to the contents shown below.

```
#ifndef lint
static char *RCSid = "$Header$";
#endif
/*
 * $Log$
 */
main()
{
    printf("hello world0");
}
```

What happened was that the *rcsit* command recognized that *sample.c* was a file containing a C program (because the file name ended in *.c*) : hence it decided to put in the header shown above which it considers to be the most appropriate header for a C program.

The *rcsit* program has stored various standard headers which are appropriate for different programming languages. The following table list the various programming languages that *rcsit* knows about and the file names that it will expect for each file type. If *rcsit* encounters a file whose name does not match any of the file types shown in the table it will assume that it is a C program.

Table 2-5: Programming Language assumed by *rcsit*

Programming Language	File Name
C Program	*.c
C Include File	*.h
Fortran	*.f
Shell Script	*.sh or *.csh
make	makefile, Makefile or *.mk
Manual entry (troff)	*.[1-9]
Pascal	*.p
TeX	*.tex

For example if your file was named *sample.sh* rather than *sample.c*, *rcsit* would insert the following header.

```
#
# $Header$
#
# $Log$
#
```

If you use non-conventional file names it is possible that *rcsit* will not be able to guess the correct header to put into your file. However, *rcsit* has a number of option flags which allow you to specify which programming language is contained in your file (and hence which header *rcsit* should use). The following table lists these option flags.

Table 2-6: Specifying Headers with *rcsit*

Command	Header Used
% <b>rcsit -c sample</b>	As for C-programs
% <b>rcsit -h sample</b>	As for C include files
% <b>rcsit -f sample</b>	As for Fortran files
% <b>rcsit -M sample</b>	As for files containing manual entries
% <b>rcsit -s sample</b>	As for shell scripts
% <b>rcsit -m sample</b>	As for makefiles
% <b>rcsit -p sample</b>	As for pascal files
% <b>rcsit -x sample</b>	As for TeX files

While the default header provided by *rcsit* will always insert appropriate identification markers into the file, it is possible that you might want to use a customized header message that you generate yourself. You can do this by creating a special *header* file containing your personalized header information. For example, Michael Cooper (the author of *rcsit*) used the following template for creating headers in his *.c* files.

```

/*
 * $Header$
 * -----
 *
 * $Source$
 * $Revision$
 * $Date$
 * $State$
 * $Author$
 * $Locker$
 *
 * -----
 *
 * Michael Cooper (mcooper@usc-oberon.arpa)
 * University Computing Services,
 * University of Southern California,
 * Los Angeles, California, 90089-0251
 * (213) 743-3469
 *
 * -----
 * $Log$
 * -----
 */

```

The *rcsit* command is not part of the normal RCS distribution; you can get it through the `comp.source.unix` newsgroup.

## Re-formatting Identification Markers

You will probably have noticed that RCS retains the identification marker within the working file. For example, the `$Revision$` marker is expanded to `$Revision: 2.1 $` and not just to `2.1`. Because it keeps the marker in the working file, RCS can recognize the identification string when you check-in the next version. If RCS eliminated the identification marker, on the next check-in it would be impossible to tell a revision number (or, for that matter, any other identification string) from "normal" text in the file. † However, retaining the original marker can cause RCS identification markers can look a little bit cluttered. This is normally no problem if the file contains source code. However, you might also use RCS to track the various versions of a published document. The appearance of the published document is crucial.

Jerry Peek (one of the authors at O'Reilly & Associates), has invented a simple macro call which can be used to parse the `$Header$` identification marker in *troff* based documentation preparation systems. This macro can be used to provide a somewhat neater appearance for the information contained in the RCS identification markers.

The following *troff* statements define his *rC* macro:

```

.de rC
.ds fI \\$2
.ds rV \\$3

```

---

† Anyone familiar with SCCS will realize that SCCS handles the expansion of identification markers in a subtly different way from RCS; SCCS replaces the original marker with the expanded text. In order to ensure that Identification markers can be recognized in newly checked-in source files, SCCS will not expand identification markers in writable copies of the file.

```
.ds dT \\$4
.ds tI \\$5
.ds aU \\$6
.ds sT \\$7
..
```

This macro definition can easily be added into your standard package of macros so that you won't have to type it in each time. This macro is designed to be used with the `$Header$` identification marker. The line which should be included in the original file is shown below.

```
.rC $Header$
```

This line is automatically expanded by RCS to:

```
.rC $Header: /usr/users/jerry/doc.ms,v 3.2 91/04/20 jerry Exp $
```

What the `rC` macro does is that it splits up the supplied parameter (the `$Header$` macro) into words and assigns each word to a different string register (a troff variable). The second word (the file name) is assigned to the string register `fI`. The third word (the version number) is assigned to the string register `rV`. The fourth word (the check-in date) is assigned to the string register `dT`, while the fifth word (the check-in time) is assigned to the string register `tI`. The sixth word (the author) is assigned to the string register `aU`. and the seventh word (the version state) is assigned to the string register `sT`.

The values of these string registers can be used at any point in the document. The following example shows how this might be done. The original document would contain the text.

```
.rC $Header$
.LP
Welcome to the SuperAppl user manual.
.LP
This is version \{rV of the document which was
checked in by \{aU at \{tI on \{dT.
The current state of this document is \{sT.
```

Version 1.2 of the document would produce the following output from the formatter:

```
Welcome to the SuperAppl user manual.
```

```
This is version 1.2 of the document which was checked in by jerry at 15:35:03 on 91/01/14.
The current state of this document is Exp.
```

However, in version 1.5 of the document, this section might read as follows.

```
Welcome to the SuperAppl user manual.
```

```
This is version 1.5 of the document which was checked in by brian at 12:03:25 on 91/02/22.
The current state of this document is Tested.
```

# 3

## Branches

"Cut is the branch that might have grown full strong; and burned is Apollo's laurel bough,  
that sometime grew within this learned man."

*Christopher Marlowe*  
*The Tragical History of Doctor Faustus, 1604*

RCS supports a concept known as *branches*, which allows you to develop more than one variant of a file. To understand the usefulness of branches, you must first understand the difference between a *variant* version of a file and the more common *revision* version.

### Variants and Revisions

Normally a new version of a file is a revision of an existing version: for example, you might revise a program to include new functions. When a new revision of a file is created, it logically supersedes the version which it revises. This makes intuitive sense because you expect that most users who previously used the old version will want the new revised version. Once the new version is available, the old version is only useful for archival purposes, bug-fixes, and so on.

A variant version implements something slightly different from the existing version. Assume that you have an existing version of a file which is suitable for use on UNIX systems. Next, you develop a new variant version which is suitable for use on MS-DOS systems. In this situation, you can't say that the MS-DOS version supersedes the UNIX version or vice-versa. Both are equally important, and equally of

interest. You don't want to treat the UNIX version as an obsolete part of the past--you will keep on working with it, alongside of the DOS version.

To see why variant versions are referred to as *branches*, you must look at a pictorial representation of the relationship between the versions in a file. We will draw what is called a *development graph*. In this graph, the different versions are interconnected by arrows indicating the flow of development work on a file. The following diagram shows a simple development graph for a file containing three revisions: †




Figure 3-1. Development Graph

All of these versions are connected by a simple linear relationship. (v3 revises v2 and v2 revises v1). Later you develop a MS-DOS version by making the appropriate changes to v3. You cannot say that the MS-DOS version supersedes v3 and hence you do not call it v4. Instead, you call it v3a to indicate that it is a variant of v3. For the same reason we do not put it on the same straight line as v1, v2 and v3 in the development graph. The following diagram shows the new development graph, reflecting the existence of v3a:




Figure 3-2. Development Graph with one Variant Version

After some time, you revise both v3 and v3a. You name the new versions v4 and v4a because they are revisions of v3 and v3a respectively. The development graph is extended as is shown below:

† In the following example we will not use RCS style version ID numbers for the versions. This is because, the rules that RCS uses to assign version ID numbers to branch versions are somewhat complex. Introducing the version naming scheme for branch versions at the same time as introducing the concept of branches would only confuse you. The rules that RCS uses to assign version ID numbers to branch versions are described later in this chapter.






Figure 3-3. Development Graph with Branches

You can see that the flow of development (as illustrated by the development graph) has been split into two separate lines or branches. The development graph has begun to look like a tree. (Perhaps we are stretching the analogy a bit, but you can see that if you introduced many new variants the development graph would begin to look like a tree.) The versions which are simple revisions of the original (i.e., v1, v2, v3 and v4) are called *trunk versions* because they are on the main line of development (the trunk of the tree). The set of revisions of a single variant is called a *branch* (i.e., v3a and v4a) because they look like a branch of the tree in the development graph.

## Numbering of Branch Versions

Earlier, we described the format of a version ID number. We told you that version ID number consists of a release number and a version number, separated by a dot, as in: *release.revision*. In fact this is not always true, it is only true for versions on the main line of development; branch version id numbers are more complex. A simple two part version identification number is not sufficient for branch versions.

In order to identify a branch version it is necessary to identify what branch it belongs to and which version of the branch it is. Therefore RCS assigns version ID numbers of the form *branch.branch\_revision*. The *branch* part identifies which branch this version belongs to. The *branch\_revision* is a single number which identifies which revision of the branch this is. The first version on any branch is *branch\_revision* 1, the second is *branch\_revision* 2, and so on.

RCS assigns ID numbers to each branch which identify where the branch fits into the development graph. RCS branch numbers have the form *branchpoint.branch\_number*, where *branchpoint* is the version ID number of the version from which development of the branch was started. It is possible for many different branches to start from the same branchpoint, hence *branch\_number* is used as a simple sequence number to identify branches when there is more than one branch descending from a single branchpoint. The first branch descended from a branchpoint is *branch\_number* 1, the second is *branch\_number* 2, and so on.

Putting all of this together we get a branch version ID number with four numbers separated by dots, as in:

```
branchpoint_release.branchpoint_version.branch_number.version
```

Since this numbering scheme isn't very easy to understand, let's look at a few examples. First look at the simple development graph illustrated below.




Figure 3-4. A File with One Branch

This file has one branch which is started from version *1.2*. The branch has an ID number of *1.2.1*, to indicate that it is the first branch descended from the branchpoint version *1.2*. The branch contains two versions. These are numbered, *1.2.1.1* and *1.2.1.2* to represent the fact that they are the first and second revisions of the *1.2.1* branch.

Now let's look at a slightly more complex example as illustrated by the development graph shown below.




Figure 3-5. A File with Three Branches

This file has three branches, *1.1.1* (the first branch that starts from the branchpoint *1.1*), *1.2.1* (the first branch that starts from the branchpoint *1.2*) and *1.2.2* (the second branch that starts from the branchpoint *1.2*). It also has a total of four branch versions, *1.1.1.1* (the first revision of the *1.1.1* branch), *1.2.1.1* (the first revision of the *1.2.1* branch), *1.2.1.2* (the second revision of the *1.2.1* branch) and *1.2.2.1* (the first revision of the *1.2.2* branch).

In the previous two examples the branchpoint versions were always on the main line of development. It is also possible to have branches which start from a branch version. † The following development graph

† Branches that start from a branch versions are a recipe for disaster. This is because, it can be quite difficult to visualize how the various versions relate to each other. I am only describing these type of branches for the sake of completeness, not because I

shows how the structure of the file might look after starting a branch from a branch version:




Figure 3-6. A Branch Starting from a Branch Version

In this file we have a version whose ID number is *1.2.1.1.1.1*. At first glance this six part ID number seems hard to de-cipher. However, this number was assigned by following the simple rules which we described earlier. The number *1.2.1.1.1.1* indicates that it is the first revision of the first branch starting from the branchpoint version *1.2.1.1*.

## Creating Branch Versions

You can create branch versions with the *ci* command in the same way that you would create any other version. RCS will decide to create a branch version, either because you explicitly tell it to do so, or because circumstances force it to so.

You can explicitly tell RCS that you would like to create a branch version by assigning a branch style version ID number with the *ci -r* option. For example the following sequence of commands will create a new branch version numbered *2.3.1.1*.

```
% co -l file
% vi file
etc.....
etc.....
% ci -r2.3.1.1 file
```

You will know from reading the previous section that a version with an ID number of *2.3.1.1* is the first revision of a branch which starts at the branchpoint *2.3*. This command will fail if there is no existing version with an ID number *2.3* to use as a starting point for the new branch.

You can check-out and lock this new branch version by using the command

---

recommend that you should create branches that start from branch versions.

```
% co -12.3.1.1 file
```

When you check-in the modified version of *file*, the new version will be numbered *2.3.1.2*. This identifies that it is a revision of version *2.3.1.1*, rather than a revision of the default version from the main line of development. In this way, development can continue on the *2.3.1* branch independently of the development work on the main line of development.

We mentioned earlier that sometimes RCS is forced to create a branch. You might well ask -- "why would RCS be forced to create a branch?". Let's look at an example.

Suppose you had a file with two versions *1.1* and *1.2*. If you execute a *co -l* command on this file, you will retrieve and lock version *1.2* (because it is the most recent version on the main line of development). When you check-in the changed copy of the working file with the *ci* command, RCS will assign the version number *1.3* to the new version. This new version will now be the most recent version on the main line of development.

Although version *1.2* is no longer the default version of the file, you could check-out and lock version *1.2* by executing the command

```
% co -11.2 file
```

When you try to check-in your changed version of the working file, RCS will consider the newly created version to be logically a successor to version *1.2* (because this is the version you locked). However, it is not possible to number this new version *1.3* because there already is a version numbered *1.3*. As an alternative, RCS will number the new version *1.2.1.1*. You can see from the following diagram that both version *1.3* and version *1.2.1.1* are logical descendants of version *1.2*.




Figure 3-7. Version 1.2 with two Successors

## Retrieving Branch Versions

You can retrieve a branch version by specifying its version id number with the *co -r* command. For example, the following command fetches the branch version of the file *sample.c* numbered *2.1.2.4*.

```
% co -r2.1.2.4 sample.c
```

If you specify a branch number rather than a particular version number, the system will fetch the most recently created version on that branch. For example, the following command fetches the the most recently created version on the *2.1.2* branch of the file *sample.c* (which might be *2.1.2.4* or *2.1.2.6*)

```
% co -r2.1.2 sample.c
```

## Setting a Default Branch

Normally, the default version fetched by each *co* command is the most recently created version on the main line of development: i.e., the line of development which has 2-digit ID number. However, if you are doing most of your work on a branch other than the main line of development, you might want to make that branch the default branch. You can set a default branch with the *rcs -b* command. Once a default branch has been set for a file, all subsequent *co* commands fetch the most recently created version from that branch, unless you explicitly ask for something else. For example, the command:

```
% rcs -b2.1.2 sample.c
```

makes the 2.1.2 branch the default branch of the file *sample.c*, and the command:

```
% co sample.c
```

fetches the most recently created version from the 2.1.2 branch. If you give the *rcs -b* command without specifying a branch number, the default branch will be reset to the main "trunk" of development. For example, the command:

```
% rcs -b sample.c
```

causes all subsequent *co sample.c* commands to fetch the most recently created version from the "trunk."

## When to use Branches

Branches allow you to maintain two parallel lines of development within one file. If you need to develop two variants of a file, you might well ask: "when should I use two branches within one file and when should I use two separate files?" It is very much up to your own personal preference. Many people like using branches because they indicate that the two branches are only slight variations of the one file. Others avoid using branches because they feel that the mechanisms are too complex and hence prone to errors. You should thoroughly familiarize yourself with the mechanisms for managing branches before deciding to use them.

Although branches were originally conceived as a mechanism for developing software for different operating systems, I personally try to avoid using branches for this purpose. This is because it can cause problems when a patch is applied to one operating system branch and not to another. If you are a C programmer, the *#ifdef* directive can normally be used to allow one source file to appear to have slightly different contents depending upon the system used to compile it. Unfortunately, not all programming languages have an equivalent feature.

However, I do find that branches are very convenient for tracking patches to an old release. For example, if version 1.5 of *sample.c* was shipped as part of release 1 of a product, work would probably start straight away to develop release 2 of the product. This would result in versions 2.1, 2.2, etc. being created as successors of version 1.5. After some time a bug report might come in which necessitates creating a patched version of *sample.c*. This patched version of version 1.5 could not be assigned a version number 1.6 since version 1.5 already has a successor. The solution would be to create a branch

starting at 1.5 which contains the patches to release 1. This would give a development graph as shown in the following diagram.




Figure 3-8. Using a Branch to Maintain an Old Release

## 4

## Symbolic Names and States

"The V sign is the symbol of the unconquerable will of the occupied territories,...."

*Winston Churchill*  
1941

If you are working on a large project you are likely to have many source files. Each of these source files will be continually be updated by many different developers. Unless you are careful you will get very confused. In this chapter we will describe how you can use the symbolic names feature and the states feature to avoid this confusion.

## States

If you are working on a large group project, it is important to know what versions of each file exist. It is also important to know how extensively each version has been tested, or what stage this version is at in the development cycle. RCS allows you to store this information by assigning a *state* to each version. A *state* is really nothing more than a short, arbitrary label that can be assigned to a version of each file. By default, RCS assigns each version of a file the state "Exp," which stands for "experimental." However, there is no intrinsic meaning to this label, or to any other. A project group can define as many states as you want, and invest them with as much (or as little) significance as it deems appropriate.

If you decide to use the state facility, your working group will need to establish some conventions about what states to use, and how they are defined. Here is one possible set of definitions:

Table 4-1: A Typical Set of States

State	Meaning	Definition
<b>Exp</b>	Experimental	A newly created version which has not been tested
<b>Test</b>	Tested	A version which has been subjected to the first level of testing
<b>Rel</b>	Released	A version which has passed QA testing and may be shipped to customers
<b>Fail</b>	Failed	A version which has failed some test

By default, RCS assigns the state "Exp" to every version when it is checked-in. By convention, this state indicates that the version is experimental and has not undergone any testing. You can override the default state for a new version by using the `ci -s` command. For example, if you subjected the new version to the first level of testing prior to checking it in, you can assign the state "Test" to this version by giving the command:

```
% ci -sTest sample.c
```

This checks-in a new version of *sample.c*, assigning the newly created version a state of "Test."

A more common (and probably more advisable) procedure is for developers to make their changes; check them in to RCS with the state "experimental" (or your equivalent); test; and, if the new version passes the testing, to change the state retroactively. You can change the a file's state after you have checked it in by using the `rcs -s` command. For example the following command changes the state of the default version of *sample.c* to "Test."

```
% rcs -sTest sample.c
```

You can also use the `rcs -s` command to change the state of a version other than the default version. To do this you must supply the version id number of the version whose state you want to change as well as the state you want to change it to. This is illustrated by the following command which changes the state of version 1.3 of *sample.c* to "Test."

```
% rcs -sTest:1.3 sample.c
```

#### NOTE

If a version fails QA testing, you should not delete this version and replace it with a version that includes the bug fix. For example if versions 1.7 of *sample.c* fails testing, the state of version 1.7 should be left at *Exp* and the fixed version should be checked-in as version 1.8.

You might be tempted to delete the failed version with the command `rcs -o1.7 sample.c`, and then check-in the fixed version as 1.7. However, this is dangerous because it introduces the possibility of confusion about "which version 1.7" is currently in use.



Given that you have defined a meaningful set of development states, and that you have diligently put "state" information into the RCS database about a file, what can you do with it? Most simply, other developers can look at your "state" information to find out what shape the project is in. Let's assume that you are using the four states that we defined earlier. It would be unwise and probably pointless to start final QA testing if one or more of the files is marked "Failed," or if one or more files are still at the "Experimental" level. But how do you read this information? If you have placed the `$State$` marker into the file (or the more elaborate `$Header$` or `$Id$` markers), the expanded version of these markers will contain the file's state. More simply, you can use the `rlog` command to dump all of the file's RCS information, including the state (plus the descriptive message, all the log messages, and so forth). We'll discuss `rlog` later.

You can also use a file's state to select which version of the file you want to check-out. The `co -s` command checks-out the most recent version of the file with a given state. For example, if you are creating a distribution tape that will be shipped to customers, you must ensure that you only check-out versions which have the state "Rel" (released; final QA passed). To do so, use the following command to check-out the most recent version of `sample.c` that has a state of "Rel":

```
% co -sRel sample.c
```

As we've said, the states that we listed above and used in our examples are only examples of how you might use RCS's "state" facility. You are free to use a version's "state" to store any information you wish, provided that everyone who uses the file is clear about the meaning of the stored information.

## Symbolic Names

If you have a large system containing many source files, it will often be difficult for you to remember which version of each source file should be used for a particular configuration. You can use symbolic names to solve this problem. Supposing you shipped a test version of our example system to selected customers and this test version was built using version 1.5 of `system.c` and version 1.2 of `system.h`. You could track this fact by assigning the symbolic name `test_release` to these versions.

Symbolic names are assigned with the `rcs -n` command, as shown below:

```
% rcs -ntest_release:1.5 system.c
% rcs -ntest_release:1.2 system.h
```

The system now considers the ascii string `"test_release"` to be equivalent to version number 1.5 in `system.c` and version number 1.2 in `system.h`. This relieves you from needing to remember which version of each source file was used for the test release. You only need to remember the name that you assigned to the release. Then the following `co` command automatically retrieves the versions you want (version 1.5 of `system.c` and version 1.2 of `system.h`):

```
% co -rttest_release system.c system.h
```

If you issued a new test release consisting of version 1.7 of `system.c` and version 1.3 of `system.h`, you might want to re-assign the symbolic name `test_release` to refer to these new version numbers. However, if you use the commands

```
% rcs -ntest_release:1.7 system.c
% rcs -ntest_release:1.3 system.h
```

the system returns an error message indicating that the symbolic name *test\_release* has already been assigned. To override the assignment of an existing symbolic name you must use the *rcs -N* command *i.e.*:

```
% rcs -Ntest_release:1.7 system.c
% rcs -Ntest_release:1.3 system.h
```

You can also delete a symbolic name without reassigning it by using the **rcs -n** command:

```
% rcs -ntest_release system.c
% rcs -ntest_release system.h
```

As a convenience, you can assign a symbolic name to a new version at the same time as creating it by using the *ci -n* or *ci -N* options. For example, the following command will assign the symbolic name "test\_release" to the newly created version of *system.c*:

```
% ci -ntest_release system.c
```

However, the above command would fail if the symbolic name "test\_release" had already been assigned. By using the **ci -N** command, we can ensure the command will succeed even if the symbolic name "test\_release" has already been assigned. For example:

```
% ci -Ntest_release system.c
```

Symbolic names are extremely useful. However, to gain maximum advantage out of symbolic names, you must be careful to adopt clear conventions about what each symbolic names means.

I find it convenient to introduce the distinction between *permanent* and *changeable* symbolic names. Strictly speaking all symbolic names are changeable because it is possible to redefine them. However, I will use the term permanent symbolic names to refer to those names that are not likely to be changed to distinguish them from changeable symbolic names whose definitions are frequently changed.

For example your local convention might be to use the following changeable symbolic names.

Table 4-2: Changeable Symbolic Names

Symbolic Name	Meaning
<i>internal</i>	This version is currently undergoing internal testing.
<i>field_test</i>	This version is currently Undergoing Field Test.
<i>released</i>	This version is the most recent version which has been released.

As part of the procedure to build a tape for use in field testing, the release engineer should assign the symbolic name *field\_test* to the appropriate version of each RCS file. Hence, if you issue the command:

```
% co -rfield_test sample.c
```

you will get a copy of the version of the file *sample.c* that is currently in field test (e.g., version 1.7). However, if you issue the same command in a years time, you will probably get a different version (e.g., version 2.11). This is because a different version of the file will be undergoing field test at that time.

If you adopt this policy for assigning symbolic names, it should help you to select the appropriate version of each source file. If you receive a bug report from a user who is using the field test version of your software, the first thing which you must do is build an equivalent version of the system on your computer. Assuming that you have assigned the symbolic name *field\_test* to the appropriate version of each RCS file, it should be simple for you to perform this build. The commands that you use would be:

```
% co -rfield_test RCS/*
% make binaries
```

If the bug report related to the released version of your software, the commands that you use would be:

```
% co -rrreleased RCS/*
% make binaries
```

The symbolic names *internal*, *field\_test* and *released* are changeable, because the association between the symbolic name and the actual version number is likely to change over time (e.g., every time a new field test is started). It is also useful to have some permanent symbolic named whose association is likely to remain fixed throughout the life of the file.

For example, a customer might phone you to say that he/she has a problem with release 2.3 of your software. However, the bug is not in version 2.3 of the source file, because release 2.3 was built using the following version 2.1 of *prog.c*, version 2.2 of *prog.h* and version 4.10 of *utils.c*.

It would be very difficult for you to remember which version of each source file was used for each release. Symbolic names help you, by providing a simple mechanism for the release engineer to record which versions are used in the release. For example the following commands define the symbolic name *two\_three* to point to the versions which were used in release 2.3 of the system.

```
% rcs -ntwo_three:2.12 prog.c
% rcs -ntwo_three:2.2 prog.h
% rcs -ntwo_three:4.10 utils.c
```

Once the symbolic name *two\_three* has been defined, you can use the following command to check-out the versions of the source files which were used with release 2.3 of the system.

```
% co -rtwo_three prog.c prog.h utils.c
```

The symbolic name *two\_three* should never be redefined, because redefining the symbolic name would erase the record of what version was used with release 2.3. Hence, *two\_three* is called a permanent symbolic name. The use of such a permanent symbolic name is subtly different from the use of a changeable symbolic name, such as *field\_test*, which is redefined every time a new field test is started.

In the preceding discussion, I assumed that the symbolic names contained complete version numbers. However, symbolic names can also be used to store either branch or release numbers.<sup>†</sup>

For example, if the branch number 1.2.1 of the file *sample.c* contains code which works on the VMS operating system, the symbolic name *VMS* can be defined to be equivalent to the 1.2.1 branch with the following command.

```
% rcs -nVMS:1.2.1 sample.c
```

Once this name has been defined, users can retrieve the most recent version from the 1.2.1 branch using the following command:

```
% co -rVMS sample.c
```

Users can also specify that they want the third revision of this branch (*i.e.*, version 1.2.1.3), with the following command:

```
% co -rVMS.3 sample.c
```

Version states and symbolic names are both used to help in selecting the correct version to retrieve with the *co* command. Hence, it is easy for you to confuse the two mechanisms. The following table clarifies the differences between them.

Table 4-3: Version States v Symbolic Names

Version States	Symbolic Names
Each version has one and only one state.	There can be one, none or many symbolic names which point to any particular version.
Several versions may have the same state.	Each symbolic name can only point to one version number

A key point to remember is that, since many versions can have the same state, it is possible to specify a version number as well as a state to the *co* command. However, since a symbolic name can only point to one version number, it makes no sense to specify a version number as well as a symbolic name to the *co* command.

For example, the following command specify that you want to check-out the most recent version from release 2 of *sample.c* which has a state of *Test*.

```
% co -r2 -sTest sample.c
```

---

RCS does not test to see that a symbolic name translates to a valid version number. Therefore, the fact that the symbolic name *sym\_name* translates to *1.3* does not guarantee that version 1.3 exists.

However, the following command is invalid because the *-rfield\_test* and *-r2* options clash.

```
% co -r2 -rfield_test sample.c
```

# 5

## Additional Utilities

"Nothing can have value without being an object of utility, ..."

*Karl Marx*  
*Capital, 1867-1883*

The primary purpose of RCS is to store multiple versions of a source file. Creating library files and checking working files in and out are the most version control operations. However, there's much more; RCS includes many additional features and utilities which help you to manage the multiple versions that you have stored. In this chapter, we will describe the most commonly used utilities that are available with RCS.

## Displaying Differences

Whenever you get a new version of a program that does not work, you inevitably ask yourself is "what is different about the new version?" If you still have a copy of the old code, you can print out both versions of the file and manually compare them. However, this process can be very tedious if the program is large and the differences are small. UNIX provides a way to automate the process: you can use the *diff* command to compare two files. For example, the following command compares the contents of *oldfile* and

*newfile*:

```
% diff oldfile newfile
```

The *diff* command prints the differences between the files as a list of the changes which must be made to *oldfile* to transform it into the *newfile*. If both files are identical, there will be no output. The *diff* command has several option flags for selecting a the format of its output. You can learn more about the various options available by looking at the *diff* command in the user's manual that came with your system.

RCS provides a special version of the *diff* command that can display the differences between different versions of files that are under RCS control. This command, called *rcsdiff*, is a convenient "front end" to the standard *diff* command. For example, the following command displays the differences between version 1.3 and 1.6 of the library file for *sample.c*

```
% rcsdiff -r1.3 -r1.6 sample.c
```

It is useful to remember that *rcsdiff* does not actually perform the comparison itself. It checks-out the relevant versions into temporary files, and then calls *diff* to do the real work. If you want to learn more about the comparison methodology, read the documentation for *diff*.

You can use *rcsdiff* to compare the current working file with a version in the library file. For example, the following command compares version 1.3 of the library file with the current working file:

```
% rcsdiff -r1.3 sample.c
```

In other words, *rcsdiff* always compares two versions of the file. If you only specify one version on the command line (only one *-r* option), *rcsdiff* assumes that you want to compare the given version with the current working file for *sample.c*. If you omit the *-r* option entirely, the system assumes that you want to compare the current working file with the default version of the library file.

The *rcsdiff* command usually issues diagnostic messages indicating which versions it is comparing: These diagnostic messages can be suppressed with the *-q* option. Any option flags aside from *-q* and *-r* are passed on to the *diff* command. For example, the following command will compare the contents of the current working file *sample* with version 1.6 of the associated library file *RCS/sample,v* using the *diff -e* command without issuing any diagnostic messages.

```
% rcsdiff -r1.6 -q -e sample
```

## Merging

Sometimes you want to make the same changes to two different files--that is, you have made a series of changes to one file, and want to apply the same changes to another, similar file. The *merge* command, which is part of the RCS system, does this automatically. Assume that you have two files *sample* and *variant*, whose contents are:

<i>sample</i>	<i>variant</i>
line1	line1
line2	variant-line2
line3	line3
line4	line4

These two files only differ in line 2. Now, assume that you make some changes to line 4 of *sample*, creating the file *sample.new* whose contents are:

```
line1
line2
line3
new-line4
```

You would like to make the same change to *variant*. Rather than manually making the change with an editor, you can use *merge* to incorporate the changes between *sample* and *sample.new* into *variant*. The command to do this is:

```
% merge variant sample sample.new
```

This command computes the changes which were made to *sample* to transform it into *sample.new*, and then incorporates these changes into *variant*. After *merge* has finished, *variant* will look like this:

```
line1
variant-line2
line3
new-line4
```

Take careful note of the order in which the files are specified. The file you want to change is listed first; the file that you have already changed by hand appears last; and the original file that you edited to create the new file is in the middle. If you specify the files in wrong order, the effect of the command will be completely different.

In our simple example, the changes between the files were so straight forward that you could have easily made the changes yourself manually. In most realistic applications, the changes to be merged will be much more complex. Sometimes, however, the changes that you wish to combine are so complex that the *merge* command can't do the job by itself. For example, let's assume that you changed line 2 of *sample.new*, in addition to line 4, so that you initially have the following three files:

<i>sample</i>	<i>variant</i>	<i>sample.new</i>
line1	line1	line1
line2	variant-line2	new-line2
line3	line3	line3
line4	line4	new-line4

Now all three files have a different version of line 2. In this case, *merge* can't tell which version of line 2 to put into the output: do you want the variant version, or the new version? This is called an overlap: the *merge* command issues an error message and flags the overlap in the output file as shown below:

```
line1
<<<<<< variant
variant-line2
=====
new-line2
>>>>>> sample.new
line3
new-line4
```

This tells you that you that the system cannot decide whether to use line 2 from *variant* or from *sample.new*. Because it cannot make an intelligent choice, it shows you both options and lets you chose which you prefer. You should manually edit the output file to select one of the alternatives.

If you use the *-p* option, *merge* displays its result on your terminal so you can "preview" them, rather than incorporating the changes directly into the *variant* file. You can also re-direct standard output so that the result goes into a different file. It is often useful, and always safer, to look at the changes that *merge* will make before incorporating them permanently into another file.

The *rcsmerge* command provides a convenient front-end for using the *merge* command with RCS library files. For example, the following command merges the changes that lead from version 1.3 to version 1.6 of *sample* into the current working file:

```
% rcsmerge -r1.3 -r1.6 sample
```

It is equivalent to the following sequence of commands

```
% co -r1.6 -p sample > sample_v1.6
% co -r1.3 -p sample > sample_v1.3
% merge sample sample_v1.3 sample_v1.6
% rm sample_v1.6
% rm sample_v1.3
```

When you need to merge different versions from a library file, *rcsmerge* is much simpler than *merge*. Note that *rcsmerge* also support the *-p* option, which lets you preview its output on your terminal.

You will often want to check-out various versions of a file and then perform a merge on the various versions. RCS provides a short-cut that allows you to save yourself some hassle by performing the merge and the check-out at the same time. To do so, use the *-j* (join) option to the *co* command. The *-j* option indicates that you would like to join (*i.e.* merge) a set of changes into the version being checked out.

For example, the flag *-jrev1:rev2* indicates that you would like to merge (join) the changes between *rev1* and *rev2* into the version that you are checking-out. For example, the *co -j* command below has the same effect as the preceding *co* and *rcsmerge* commands:

```
% co -r1.4.2 sample.c
% rcsmerge -r1.4.1.3 -r1.4.1.4 sample.c
```



```
% co -r1.4.2 -j1.4.1.3:1.4.1.4 sample.c
```

If you only specify one version number with the `-j` option, the system looks for the version which is the common ancestor of the version you are checking-out and the version you specified. It then merges the changes that have taken place between the constant ancestor and the `-j` file into the version indicated by the `-r` option (or the default version, if `-r` isn't given). For example, the following command merges the changes between version 1.2 and version 1.2.1.3 into the copy of version 1.4 which is being checked-out:

```
% co -r1.4 -j1.2.1.3 sample.c
```

In this case, version 1.2 is the common ancestor of version 1.4 and version 1.2.1.3. That is, 1.2 is the "branchpoint" at which branch 1.2.1 diverges from the main tree, to which 1.4 belongs. Note that the "common ancestor" will always be a branchpoint.

You normally use the `-j` option to merge additional changes into the working file. However, you can also use `-j` to exclude certain changes from the working file. For example, consider the command:

```
% co -r1.2 -j1.3:1.5 sample.c
```

This command checks-out a working file that contains version 1.2 of *sample.c* merged with the changes between 1.3 and 1.5. This is effectively the same as checking-out a copy of version 1.5, except that the changes between 1.2 and 1.3 are not included. If you decide that version 1.3 was a complete mistake, but you are still optimistic about some code that has been added in versions 1.4 and 1.5, you might find this command useful. Why might this command be useful? Let's assume that version 1.3 added one feature which just never worked the way you had planned, but 1.4 and 1.5 added substantial new features that haven't been a problem. Rather than starting from 1.2 and copying the features from 1.4 and 1.5 into the new file, you can use this command to automatically delete the features you added for version 1.3.

## Deleting Old Versions

RCS library files are designed to minimize the space needed for storing many different versions of a file. However, they do grow as you check-in additional versions. From time to time, you will want to delete some of the old, less-useful versions from a library file in order to save room. You can do this by using the `rcs -o` command. For example, the following command deletes version 1.3 of *sample.c*:

```
% rcs -o1.3 sample.c
```

You can delete more than one version with the same command. There are three ways of specifying the range of versions which you wish to delete. First, you can use the notation `-orev1-rev2`: this tells RCS to delete all versions between *rev1* and *rev2*, including *rev1* and *rev2*. For example the following command will delete versions 1.2, 1.3 and 1.4 of *sample.c*.

```
% rcs -o1.2-1.4 sample.c
```

The version range can include branch versions, so long as both *rev1* and *rev2* belong to the same branch (or the trunk) of the development tree. For example, the following command will delete branch versions 2.2.1.2, 2.2.1.3 and 2.2.1.4 of *RCS/sample.c,v*.

```
% rcs -o2.2.1.2-2.2.1.4 sample.c
```

The second way to specify a range of versions is to use the notation `-orev-`. This means to delete all versions from *rev* onwards. If *rev* is on the main line of development, *rev* and any more recent version on the main line of development will be deleted. If *rev* is a branch version, *rev* and any more recent version on that branch will be deleted. For example, the following command will delete all versions of *sample.c* from version 2.4 onwards.

```
% rcs -o2.4- sample.c
```

The following command will delete all versions on the *1.2.1* branch from version *1.2.1.3* onwards (e.g. versions *1.2.1.3*, *1.2.1.4*, *1.2.1.5*, etc.).

```
% rcs -o1.2.1.3- sample.c
```

The third way, you can express a range of revisions is by using the notation `-o-rev`. If *rev* is on the main line of development, this means "delete all versions on the main line of development up to and including *rev*." If *rev* is a branch version, this means "delete all versions on the branch, up to and including *rev*." For example, the following command will delete all versions of *RCS/sample.c,v* up to and including version *1.4* (e.g. versions *1.1*, *1.2*, *1.3* and *1.4*).

```
% rcs -o-1.4 sample.c
```

Likewise, the following command will delete all versions of the *1.2.2* branch up to and including version *1.2.1.4* (e.g. versions *1.2.1.1*, *1.2.1.2*, *1.2.1.3* and *1.2.1.4*).

```
% rcs -o-1.2.1.4 sample.c
```

You cannot delete any version which is locked, nor can you delete a version that is a branchpoint (i.e., the starting point for a branch). If you want to delete a locked version, you must first remove the lock with the `rcs -u` command. If you want to delete a branchpoint version, you must first delete all of the versions on the branch with a separate `rcs -o` command.

## The *rcsclean* command

If you are a typical computer user, you frequently come under pressure to reduce your disk space. At these times, you scan all of your directories to see what information (if any) you have stored redundantly. You can safely delete any working files that have not been modified since they were checked-out--you can always re-create them by checking-out another working file from the RCS file. The *rcsclean* command does this automatically. It eliminates the need for examining your working files by hand to see if they can be safely deleted. For example, both of the following commands:

```
% rcsclean sample.c
--and--
% rcsclean RCS/sample.c,v
```

compare the contents of the working file *sample.c* with the contents of the default version of the RCS file *RCS/sample.c,v*. If they are identical, *rcsclean* deletes the working file *sample.c*. If you have locked the default version of *RCS/sample.c,v*, *rcsclean* also deletes the lock. If the working file and the default

version from the library file are not identical, *rcsclean* does nothing.

Note that *rcsclean* doesn't do everything that you'd like it to. It only compares your working file with the default (usually, the most recent) revision. Because you can reconstruct any earlier revision from the library file, it is safe to delete the working file if it is identical to any of the versions of the RCS file. However, it isn't practical to compare the working file with each older version. If you want to check your working file against a particular version, you can use the *-r* option to *rcsclean*, which selects an alternative version for comparison with the working file.

For example the following command compares the contents of the working file *sample.c* with the contents of version 1.3 of the RCS file *RCS/sample.c,v* (and deletes the working file if they are identical).

```
% rsclean -r1.3 sample.c
```

Likewise the next command compares the contents of the working file *sample.c* with the contents of the most recent version of release 2 of the RCS file *RCS/sample.c,v*.

```
% rsclean -r2 sample.c
```

The *rcsclean* command would not be terribly useful if it only compared one file at a time. Fortunately, it can do much more. The *rcsclean* command is commonly used to check a group of files. For example the following command compares every file ending in *.c* with the default version of the corresponding RCS file (and deletes the working file if they are identical).

```
% rsclean *.c
```

Likewise the next command compares the default version of every RCS file in the *RCS* directory with the corresponding working file.

```
% rsclean RCS/*,v
```

The *rcsclean* command, unlike most UNIX commands, is fairly verbose. It always tells you which versions of each file it is comparing and what files (if any) it is deleting. Therefore, you should never be in doubt about its effects.

### NOTE

The *rcsclean* command is a relatively new addition to RCS. If you are using an old release of RCS the command might not be available.

## The *rcsfreeze* command

In an earlier chapter, we described how you can use symbolic names to record which version of each individual file was used to build some version of the system as a whole. You will frequently need to assign a symbolic name to "the default version" of all of the RCS files. For example, when your system officially enters beta test, you'll want to assign the name *beta\_test* to the current version of each of your files. However, it isn't really convenient to do this with *rcs -n*. The *rcs -n* command requires you to state an explicit version number whenever you associate a version with a symbolic name. This means

that you must first find out the version number of the default version of each file. Then, you will probably need to type a long sequence of commands like the following:

```
% rcs -nbeta_test:2.7 prog.c
% rcs -nbeta_test:2.2 prog.h
% rcs -nbeta_test:2.5 routines.c
% rcs -nbeta_test:2.1 routines.h
```

To simplify this procedure, RCS provides a special command called *rcsfreeze*. It automates the task of assigning a symbolic name to the default version of all the RCS files in the current directory. It finds all the RCS files in the current directory (or the *RCS* sub-directory if it exists), determines the version number of the default version stored in that file, and assigns a symbolic name to that version number. For example, the following command assigns the symbolic name **beta\_test** to the default version of all the RCS files:

```
% rcsfreeze beta_test
```

Note that *rcsfreeze* invokes the *rcs -n* command rather than the *rcs -N* command. This means that it cannot re-define a symbolic name that is already in use. The assignment will fail if the symbolic name **beta\_test** has already been assigned to any version of any RCS file.

If you do not specify a symbolic name (i.e., if you invoke *rcsfreeze* with no arguments), the system chooses a name for you. This name has the form "C\_x," where **x** is a sequence number that is incremented whenever **rcsfreeze** is run. *Rcsfreeze* tells you the sequence number when it is assigned. It also records the sequence number in a file named *.rcsfreeze.version* so that it can ensure sequence numbers are not repeated.

When you give the *rcsfreeze* command, the system prompts you for a log message describing why you assigned the symbolic name. It then stores the log message (along with the symbolic name, the sequence number, and the current date) in a file called *.rcsfreeze.log* for future reference. The *.rcsfreeze.log* file contains entries like:

```
Version: release_1(C_2), Date: Tue Sep 18 15:31:12 WET DST 1990
-----
    creating the first release to be shipped to customers
    this will have the symbolic name release_1
-----
Version: beta_test(C_1), Date: Tue Sep 18 15:27:32 WET DST 1990
-----
    Creating a beta test version of the entire system
-----
```

You should beware that *rcsfreeze* only assigns symbolic names to versions that have been checked-in. If you have made changes to any of the files since they were last checked-in, you must perform a check-in before you run *rcsfreeze*. For this reason, many users run *rcsclean* immediately prior to *rcsfreeze*. If everything is properly checked-in, *rcsclean* will delete any working files that are remaining. Once you know that everything is checked-in, you can run *rcsfreeze* without worrying.

#### NOTE

The *rcsfreeze* command was developed by Stephan Bechtolsheim. It was not initially part of RCS, and may not be part of your system. If *rcsfreeze* is not available on your system, you can get a copy from the Free Software Foundation.

The *rlog* command

Now, let's consider another problem. After returning from vacation, you go to your desk and compile your file *sample.c*. It doesn't work--or it works, but its behavior seems to be slightly different. Strange--you're sure that everything was fine when you left. In order to figure out what went wrong, you must answer a number of questions about *sample.c* file, including:

- Have any new versions of *sample.c* been created in your absence?
- Who created these new versions?
- Why did they create these new versions?

The *rlog* command answers these questions and more by displaying all the information that the system has stored about the file. To ask the system to print out everything it knows about the file *sample.c*, use the following command:

```
% rlog sample.c
```

The output of this command might be:

```
RCS file:      RCS/sample.c,v;   Working file:   sample.c
head:         1.3
branch:
locks:        mary: 1.3; fred: 1.2;  strict
access list:   fred mary  harry
symbolic names: middle_rev: 1.2;
comment leader: " * "
total revisions: 3;  selected revisions: 3
description:
sample file used to illustrate the use of the rlog command
-----
revision 1.3          locked by: mary;
date: 89/08/07 17:07:18;  author: mary;  state: Exp;  lines added/del: 31/14
Altered the contents of the file significantly
-----
revision 1.2
date: 89/08/07 17:06:34;  author: fred;  state: Exp;  lines added/del: 3/0
Inserted some additional lines
-----
revision 1.1
date: 89/08/07 17:04:27;  author: mary;  state: Exp;
Initial revision
=====
```

If you've followed our discussion all along, this output should be self-explanatory. However, here's a summary of what the system is telling us:

- The name of the RCS file is *RCS/sample.c,v* and the name of the working file is *sample.c*.
- The most recently created version (called the *head* version) is numbered 1.3.
- No default branch has been set--i.e., by default, check-outs and check-ins are to the "main" branch, or "trunk."
- Version 1.3 is locked by *mary*. Version 1.2 is locked by *fred*. Strict locking is enforced. When strict locking is not enforced, the word "strict" does not appear at the end of the list of locks.

- The access control list contains the names *fred*, *mary* and *harry*. These are the only users allowed to make modifications to the RCSfile.
- The symbolic name "middle\_rev" is equivalent to the version number 1.2.
- The comment leader is " \* ". This means that every line of the the expanded \$Log\$ marker will be preceded by the string " \* ".
- This RCS library file contains three versions of *sample.c*. This *rlog* command is going to give us detailed information about all three versions. Later, we will see how to reduce the volume of output from *rlog* by restricting the versions that it will report on.
- The descriptive text stored in the RCS file is: "sample file used to illustrate the use of the rlog command."
- The first version reported upon is version 1.3: It is locked by "mary," who is also its author.
- It was checked-in on August 7, 1989 at 7 minutes and 18 seconds past 5pm. Note that the date is printed as year/month/day rather than month/day/year. RCS uses this notation because it is equivalent to the "most significant bit first" format.
- The state of the version is "Exp," which usually (but not necessarily) means "experimental."
- The changes made to the previous version (version 1.2) to create this version involved deleting 14 existing lines and inserting 31 new lines.
- When "mary" checked-in this version, she gave the log message: "altered the contents of the file significantly"
- Similar information is provided about the other two versions 1.2 and 1.1; we won't bother to summarize it.

If you're "fred", and you left for vacation just after checking-in version 1.2, you can now start investigating why the program is different. Rather than puzzling over the code and trying to remember what changed, you can go directly to the author of those changes, Mary, and ask her what she did while you were gone.

The name *rlog* is an abbreviation for "read log messages." However, as we saw in the example, the *rlog* command prints much more than just the log messages. The output from *rlog* can be very long, even for a simple file like the one we illustrated. If you don't want to read the entire report, there are many alternatives--perhaps too many. *rlog* has an almost interminable list of options designed to customize the report in one way or another. First, the *-t* option gives general information about the file under RCS control, but omits all of the information about specific versions. The command

```
% rlog -t sample.c
```

prints the same information as the default *rlog* command, up to and including the descriptive text (the first 11 lines in our example), but omits everything about the individual versions. You can reduce the amount of output even further by using the *rlog -h* option. The command *rlog -h* produces the same output as *rlog -t*, except that the descriptive text will not be printed. In other words, it prints everything up to, and including, the "total revisions" line.

While you probably don't want to see a long list of details about all the versions in the file, you might be interested to find out about some of them. You can specify which versions you care about by using the *-r* option. For example, the command:

```
% rlog -r1.2 sample.c
```

prints the usual header information but only gives details about version 1.2 of *sample.c*. If you specify a

branch number in the place of a version number, the system prints details for all versions on that branch. For example, the command:

```
% rlog -r1.2.1.1 sample.c
```

reports upon all the versions on branch 1.2.1 of the file *sample.c*.

You can request details about several version by specifying a list of versions. For example the following command prints version-specific data for versions 1.2 and 1.4 of *sample.c*.

```
% rlog -r1.2 -r1.4 sample.c
```

As well as specifically naming the versions you want details about, you can also specify a range of versions. For example the next command prints version-specific data for versions *sample.c* in the range 1.2 to 1.5 (i.e., 1.2, 1.3, 1.4 and 1.5).

```
% rlog -r1.2-1.5 sample.c
```

Likewise, the next two commands will print information about versions of *sample.c* up to and including 1.3 (i.e., 1.1, 1.2 and 1.3) and from version 1.3 onwards (e.g. 1.3, 1.4 and 1.5).

```
% rlog -r-1.3 sample.c
```

```
% rlog -r1.3- sample.c
```

You can also specify which versions you want by giving a check-in date. To do so, use the command *rlog -d*. RCS allows you to specify dates in free format as described in appendix D. You can specify a range of check-in dates by using the notation *start\_date<end\_date*. For example, the command:

```
% rlog -d'1-AUG<10-SEPT' sample.c
```

reports details about all versions of the file *sample.c* that were checked-in between the dates August 1 and September 10.

You can also print details about all versions checked-in before or after a certain date. For example the following two commands will report on all versions checked-in before the first of September and after August 15 respectively.

```
% rlog -d'<1-SEPT' sample.c
```

```
% rlog -d'>15-AUG' sample.c
```

Sometimes you might wish to report on versions that were checked-in during any one of a number of different date ranges. You can do this by listing all of the date ranges with the *-d* option, using a semi-colon to separate them. For example, the command:

```
% rlog -d'10:30<17:00;3-SEPT>1-SEPT;<1-AUG' sample.c
```

prints information about all versions of *sample.c* that were checked in between half past ten this morning and five o'clock this evening, as well as versions that were checked in between the first and third of September or before the first of August.

## NOTE

It is extremely important to put single quotation marks around the date range. Otherwise the shell will interpret the '<' and '>' characters as indirection requests for standard I/O, and the result would be significantly different from what was intended. Likewise, if you don't use quotation marks the shell will interpret ';' as a command delimiter.

The `-w` option restricts reporting to versions which were created by a particular author. For example, the following command only reports on versions that were checked-in by "fred."

```
% rlog -wfred sample.c
```

Likewise the following command reports on versions that were checked-in by "fred," "mary" or "harry."

```
% rlog -wfred,mary,harry sample.c
```

You might like to restrict your report to versions which are currently being worked upon (*i.e.* versions which are locked). The `-l` option restricts reporting to versions which are locked. For example, the following command reports on versions of *sample.c* which are locked (by anyone).

```
% rlog -l sample.c
```

Sometimes, you might like to restrict your report to versions which are currently being worked upon by particular people. If usernames are supplied with the `-l` option, *rlog* will only report on versions which are locked by the named users. For example, the following command reports on versions of *sample.c* which are locked by either "fred" or "mary."

```
% rlog -lfred,mary sample.c
```

Another way of restricting the number of versions reported upon, is to specify that only versions with a particular state should be reported upon. The *rlog* `-s` option is provided to do this. For example, the following command reports on versions of *sample.c* whose state is equal to "Test."

```
% rlog -sTest sample.c
```

Likewise, this next command reports on versions of *sample.c* whose state is either "Test" or "Rel."

```
% rlog -sTest,Rel sample.c
```

The *rlog* `-b` command can be used to restrict the report to versions on a particular branch. If *sample.c* has a default branch set, the next command prints details about versions which are on the default branch; otherwise it prints details about versions on the branch with the highest branch number.

```
% rlog -b sample.c
```

You can combine two or more selection criterion within one *rlog* command. For example, the command:

```
% rlog -wfred,mary -sTest,Rel sample.c
```

prints information about *sample.c*, giving details about versions that were checked-in by either "fred" or



"mary" and that have a state of "Test" or "Rel."

The *rlog* command is frequently in shell scripts to convert the name of the working file to the name of the corresponding RCS file. So that you don't need to write a special shell script just to extract the RCS library file name from the rest of *rlog*'s output, the *-R* option to *rlog* prints the RCS file name only. For example, the following command prints the name of the RCS file associated with *sample.c*:

```
% rlog -R sample.c
RCS/sample.c,v
```

This could be used as part of a shell command. For example, assume that the shell variable *file* contains the name of a working file. Then the following series of C-shell commands set the shell variable *RCS\_file* equal to the name of the corresponding library file:

```
% echo $file
sample.c
% set RCS_file = `rlog -R $file`
% echo $RCS_file
RCS/sample.c,v
```

Finally, assume that you are involved in a team project with many source files, all of which are under RCS control. You might want to find out which of the files are currently under development--i.e., which files have been checked-out and locked. (Files which are checked out but not locked aren't "under development," because they can't be modified). To find out which files are locked, give the command *rlog -L* for the files that are in your working directory. The *rlog* command then prints its standard report for every file that is locked, ignoring those that aren't.

For example the following command prints a full report on all RCS files that have been locked.

```
% rlog -L RCS/*,v
```

This next command works the same as the previous, but makes its inquiries on the basis of the working files in the current directory.

```
% rlog -L *
```

The *-L* option can be combined with any other command line option. For example, the following command combined the *-L* and *-R* options to print just the names of the RCS files which have been locked.

```
% rlog -L -R RCS/*,v
```

## Telling White Lies

As well as storing the contents of each new version the *ci* command records certain housekeeping information about each new version that is created. For example, it will store who created the new version and when the version was created. RCS can find out this information itself, however, there are occasions when you might want to tell a few white lies about this housekeeping information.

Normally, the *ci* command records the date on which you execute the command as the creation date of the new version. There are times when you this isn't what you want. Suppose you created a new version of *sample.c* on the 21st of January 1990, but forgot to check it in until the February 1st. You would still like the RCS system to record the 21st of January 1990 as the creation date; someone who wants to reconstruct the system as it was on January 22 presumably wants to see your version of *sample.c*, even though it wasn't officially part of the system until much later.

The *ci -d* command lets you assign a creation date to the new version. For example, the command below checks-in *sample.c*, forcibly setting the creation date to the 21st of January, rather than the current date:

```
% ci "-d21 JAN 1990" sample.c
```

You will notice that we enclosed the *-d* option inside of quotation marks: this tells the shell that the date string is a single parameter, rather than a series of arguments separated by spaces. The date specified with the *ci -d* command can be specified in "free format" as described in appendix D.

The *ci* command normally assumes that the user performing the check-in is the version's "author." Again, there are times when you might want to do otherwise: for example, you may be checking-in a file because someone left for vacation and forgot to do his own check-in. The *ci -w* command lets you assign authorship of the new version to another user. For example, if you wanted to check-in a new version of the file *sample.c* on behalf of your friend *sally*, use the command:

```
% ci -wsally sample.c
```

If you are using RCS to store versions of software that is being developed elsewhere, it is desirable that you always should use the same version numbers as the developers. For example, assume that Jim Bean & Co. has purchased a source license for an accounting package called *Cook\_the\_Books*. This package is supplied by Dodgy Software Ltd (DSL for short). DSL use the RCS package to maintain multiple version of the source code. Every time they make a significant change to the source code, they ship a copy of the new source code to Jim Bean.

Jim Bean, whose login-id is *jim*, acts as a system manager for Jim Bean & Co. and installs each new version as he receives it. Experience has taught Jim that he sometimes needs to restore old versions. Therefore, whenever he receives a new version, he doesn't delete the old code. Instead, he maintains on-line copies of all the versions using RCS. In order for Jim to communicate with DSL effectively in relation to bugs in the software, it is desirable that Jim Bean & Co. should use the same version numbers as DSL and agree about the release date of each version.

Now, assume that a user named *fred*, who works for DSL, developed version 1.5 of the file *accounting.c* and checked it into the RCS file on January 1st. This was subjected to rigorous testing and hence its state was set to "tested." The *\$Id\$* marker was included in the original file for identification purposes. In version 1.5 of the source code, the *\$Id\$* marker is expanded as shown below:

```
$Id: accounting.c,v 1.5 90/01/20 11:26:48 fred tested $
```

Jim Bean & Co signed a license agreement on the 20th of February. Since version 1.5 had not been superseded in the meantime, it was sent to Jim on the distribution tape. Now consider what happens when Jim checks in his new source code. If he uses the simple command

```
% ci accounting.c
```

the version number would be set to 1.1, the author would be identified to be *jim*, the creation date would

be 20th of February and the state would be "Exp". It is clear that confusion arise when Jim calls to complain about a bug in version 1.1 of his software. DSL would have their own idea about what constituted version 1.1, which would be completely different from Jim's.

Jim could ensure that he uses the same version number, creation date, author and state as DSL by using the following command to check-in the source files.

```
% ci -r1.2 "-d20-JAN-1990 11:26:48" -wfred -stested sample.c
```

Before Jim could know the correct command line options to use with the *ci* command, he would have to know the correct version number, creation date, state and author. Since DSL have been considerate enough to place *id* markers in their source files, he can find out this information by looking for the *id* marker and then using its contents. Luckily, Jim does not have to go to the trouble of doing this himself, because RCS has a command which will automatically search for and use the information from the *id* markers.

The *ci -k* command tells the system derive check-in parameters from the expanded markers in the working file itself. The *ci* command searches the file for expanded identification markers. It then uses the version number, author and creation date from these expanded ID markers (if any) to provide the version number, author and creation date of the new version.

For example, if Jim used the command

```
% ci -k accounting.c
```

to check-in his files, RCS would search through the working copy of *accounting.c* and find the expanded *\$Id\$* marker. From the expanded marker, the system would know that this version was authored by *fred*, was created on 1st of January, has a state of "tested" and a version number of 1.5. RCS will then enter the file into its library with this version number, author, and date. Now, when Jim calls DSL to complain, his "version 1.5" will match theirs. Using *ci -k* can therefore eliminate much of the potential confusion.

If you use the *ci -k* command and no identifying keywords are found in the working file, the normal default values will be used for the author, creation date, version number and state. Any of the values read from the keywords in the file may also be overridden using the *-w*, *-d*, *-r* and *-s* options.

It might be important for you to be able to distinguish which versions were created with the *ci -k* command and which versions were created with the normal *ci* command. In order to assist you, RCS stored a special log message with every version that is checked in using the *-k* option. Whenever you check-in a file using *ci -k*, the system will automatically generate a log message of the form "checked in with *-k* by *user* at *date*"; where the *date* is given in the internal format used by RCS, which is explained in appendix D. For example:

```
checked in with -k by jim at 90.02.20.09.11.36.
```

# 6

## Interaction with Make

"And God said, let us make man in our image, after our likeness"

*The Holy Bible: Genesis*

*Make* is one of the UNIX system's most exciting and interesting utilities. It is commonly used to build executable systems from the component source files, and can be put to many other uses. If you have used UNIX to any extent, you are probably familiar with *make* already. Why describe *make* in a book about RCS ?

Despite the fact that *make* is used frequently in conjunction with RCS, the interface between them is far from optimal. There are a number of features of *make* which make it difficult to use with any version control system. In this section, we describe a primitive method of interfacing *make* to a version control system. We also describe the problems inherent in this interface. We won't describe *make* itself in any detail. The *make* system is too complex to be adequately described in this small section, so we assume that you are already familiar with its basic features. If you want to learn more about *make*, read your system's documentation or *Managing Projects with Make* written by Steve Talbot (O'Reilly & Associates).

Normally, *make* is used with the most recent version of the source files to build an up-to-date copy of the executable system. Therefore, by combining *make* with a version control system, you should be able to check out and build an arbitrary version of the program automatically. To see how this might happen, let's start with a simple case.

The dependencies between the various compiled files and source files are stored in a *makefile*, along with the commands required to rebuild the system. The *makefile* contains a rule for each *target* file that it knows how to build. The rule for each target file must state the *dependencies* (a list of files that are required to build this target) and a *build command* (the shell command or commands required to generate the target file from its dependencies). Here's the format of a typical rule:

```
target: dependencies
      build command
```

For example, let's look at a simple system consisting of two source files *sample.c* and *sample.h*, which produce an executable file *sample*. The file *sample.c* contains the C code for the system; it uses the header definitions contained in the header file *sample.h*. The executable file *sample* can be generated by compiling the source file *sample.c*. The Makefile for this system would be as follows

```
sample: sample.c sample.h
      cc -o sample sample.c
```

These lines tell MAKE that the executable file *sample* depends upon both *sample.c* and *sample.h*: this implies that *sample* must be recompiled if either *sample.c* or *sample.h* changes. The command to compile *sample* is **cc -o sample sample.c**.

In order to interface *make* to a source control system, you must include extra lines indicating that the source files are dependent upon the corresponding library files. You must also tell the system how to create the source file from the library file by giving the appropriate check-out command. For example:

```
sample.c: RCS/sample.c,v
      co sample.c
```

These lines tell the system that the source files can be generated by checking-out a working file from the relevant library file. They also tell the system that if the modification time of the library file is more recent than the modification date on the working file, it should check-out a new working file. With this additional dependency, *make* will check-out a new working file if it needs to get an up to date copy—for example, if a co-worker has checked-in a new version of *sample.c* since the last time you checked it out.

This interface to the source control system is adequate but not really very powerful. The main problem with this interface is that the system will always retrieve the current default version of the library file. Hence, it will not exploit the true benefit of using a version control system, which is the ability to re-create old versions. Fortunately, this problem is easy to fix. The options you provide with the check-out command determine the version of the source file that will be checked-out, and consequently the version of the executable that will be built. If we can teach *make* to provide the checkout options that specify the particular version that interests us, our job will be done. The simplest way to do this is by using MAKE variables:

```
sample.c: RCS/sample.c,v
      co $(COFLAGS) sample.c
```

These lines do tell *make* to use the contents of the *COFLAGS* variable as an option flag to the checkout command. You can define the *COFLAGS* variable on the *make* command line. For example, consider the command:

```
% make sample COFLAGS=-r2
```

When you give this command, *make* does whatever is necessary to generate the file named *sample*, using the option flag **-r2** whenever checking-out a file from the version control system. This means that the

most recent version from release 2 of each file will be checked-out and used for the compilation. By altering the value of the *COFLAGS* variable, we can vary the options used with the check-out command and hence the version of *sample* which will be generated.

More profitably, you can use the *COFLAGS* variable to supply any check-out option you wish. For example, assume that you are using RCS and that you have assigned the symbolic name **final\_release\_1.0** to the version of your product that was actually shipped sometime last year. When a bug report comes in, you can reproduce the executable that you shipped with the command:

```
% make sample VERS=-rfinal_release_1.0
```

*Make* issues the *co* commands to check-out the appropriate version of every file, regardless of the current version. You could use similar variations of the *make* command to create your product as it stood on any particular date, and so on.

## Problems

The example shown accomplishes a simple but effective interface between RCS and *make*. With this interface, *make* can use RCS to access old versions of the source files and hence build old versions of the executable files. However, there are still several problems with this interface. We will list here some of the ways in which *make* will not interact with RCS in the way that you might intuitively expect. † If you intend to use *make* a version control system, you should read and understand these oddities so that you do not get any unwelcome surprises.

- You must explicitly enter details into your *makefile* about every source file which is under version control. This can become tedious if you are using a large number of source files. Some (but by no means all) variants of MAKE have default rules which specify how source files should be checked-out from RCS. You should familiarize yourself with what default rules (if any) your system has for interacting with RCS.
- Many variants of *make* assume that all the required files are in the same directory. If you are using such a variant of *make*, it is almost impossible to specify good default rules because the library files are normally in a different directory from the source files.
- *Make* assumes that the compilation rules do not vary. However, the commands needed to compile release 2 of the system might differ from the commands used to compile release 1 of the system--particularly if you are working with a complex optimizing compiler. There is no simple way to store two alternate sets of compilation rules within one Makefile.
- We used the simple rule that a new working file should be checked-out if the current working file is older than the library file. This is not adequate for every situation. If the working directory currently contains the current version (release 4) of the source files, and we wish to create a new version of the executable using release 2 of the source files, MAKE cannot tell by looking at the modification times of the file that a new working file (release 2 of the source code) needs to be checked-out. The only solution to this problem is to force *make* to check-out a new working file. This can result in

---

When I refer to the *MAKE utility*, I mean the original MAKE utility as developed by Stu Feldman. There are several enhanced MAKE utilities currently in existence. Some of these enhanced MAKE utilities have solved some of the problems discussed in this section.

The MAKE utility supplied by GNU has advanced features which help you interface between MAKE and RCS. I will describe how these advanced features can be utilized (if you have them). It is possible that the MAKE utility, which was supplied with your variant of the UNIX operating system, has some of these advanced features.

unnecessary compilations and is not satisfactory for large systems.

We have not described some of the more subtle problems that arise when using *make* in conjunction with a source control system. Recently, some users have written enhanced versions of MAKE which provide a cleaner interfaces to RCS. There are also a number of utilities available which assist in some of the more mechanical aspects of using *make* in this context. These utilities will be discussed later in the book. All users should at least be aware that there are serious pitfalls in the interface between RCS and MAKE.

## Homegrown Solutions

Since the standard MAKE utility does not contain a good interface telling it how to check-out files from RCS, the obvious solution is to define your own interface. In this section we show you how you can improve the interface between RCS and MAKE.

Your first task should be to define a good default rule telling MAKE how to check-out files from RCS. As an example you can look at the default rule used by the enhanced GNU MAKE. *i.e.*

```
CO = /usr/new/co
COFLAGS =

%:: %,v
    $(CO) $(COFLAGS) $< $@

%:: RCS/%,v
    $(CO) $(COFLAGS) $< $@
```

These two rules use the pattern matching feature of GNU MAKE which is not available with the standard MAKE utility. The first rule says that a source file can be built from a file with the same name except with a ,v appended to the end of it. The second rule says that a source file can be built from a file with the same name except with a ,v appended to the end of it and in a sub-directory named RCS. For either rule the default action is to check-out the file using the */usr/new/co* program and with no options specified. However, an alternate check-out program can be specified in the *CO* macro and option flags can be specified in the *COFLAGS* macro. For example the following command will build the *install* target but will use the */usr/local/co* command with the *-r2* option flags whenever it needs to check-out a source file from RCS.

```
% make install CO=/usr/local/co COFLAGS=-r2
```

If your version of the MAKE utility does not support pattern matching rules, you will not be able to define such simple and elegant generic rules for checking-out source files from RCS. The standard MAKE utility does have any way of defining generic rules for fetching files from another directory. Therefore, if you store your RCS files in a sub-directory named *RCS* (as most people do), you must define explicit rules for how to check-out each of the source files from RCS. For example, your *makefile* might contain the following rules:

```
CO = /usr/new/co
COFLAGS =

prog.c: RCS/prog.c,v
    $(CO) $(COFLAGS) RCS/prog.c,v prog.c
```

```

prog.h: RCS/prog.h,v
    $(CO) $(COFLAGS) RCS/prog.h,v prog.h

utils.c: RCS/utils.c,v
    $(CO) $(COFLAGS) RCS/utils.c,v utils.c

```

If you have a large number of source files, the list of rules can become quite long. While the standard MAKE utility does not provide pattern matching rules, it does provide an almost equivalent feature called suffix rules. The suffix rules feature allows you to define how files with one suffix can be generated from files with another suffix<sup>‡</sup>. For example the following rule tells MAKE how files with a suffix of ".c" can be generated from files with the same name except for a suffix of ".c,v" (the corresponding RCS file).

```

.SUFFIXES : .c .c,v

.c,v.c :
    $(CO) $(COFLAGS) $< $@

```

The suffix rules feature is not as flexible as the pattern matching feature. An obvious problem with the above rule is that it only relates to checking out source files whose name ends in ".c". You would have to define a separate suffix rule to tell MAKE how to check-out files whose names end in ".h" and with any other file name ending that you use. Another more serious problem with this rule is that it will not look for the RCS files in any directory other than the current working directory. The suffix rules feature does not allow you to specify that the dependency files are in another directory (*e.g.*, the *RCS* sub-directory).

A simplistic solution to this problem would be to define a default rule which would check-out files from RCS. This rule would be defined as follows:

```

CO = /usr/new/co
COFLAGS =

$(CO) $(COFLAGS) $<

```

This rule tell MAKE that, whenever it has no explicit or suffix rules telling it how it can build a particular file, it should try checking-out the file from RCS. The *co* command itself will automatically look in the *RCS* sub-directory to find the RCS file. However, the problem with using a *.DEFAULT* rule is that it is only activated if a source file does not exist; the MAKE utility will not recognize that a source file needs to be checked-out again when the RCS file is updated, because the *.DEFAULT* rule does not tell it how to find the relevant RCS file to test its modification date.

The default rule provided with GNU MAKE is substantially more elegant than any default rule that you can define with standard MAKE. However, even the GNU MAKE default rule is far from perfect. The GNU MAKE rule will specify that a source file needs to be checked-out if the modification date of the RCS file is later than the modification date of the source file. Sometimes this can result in the source file being unnecessarily checked-out. For example, if a new user was added to the access control list of the RCS file, the modification time of the RCS file would be updated and MAKE would decide to perform a

---

<sup>‡</sup>In this context the file suffix is the part of the file name that follows the "."; you should consult your systems' documentation for more information about the suffix rules feature.



check-out (and recompile the source file) even though the contents of the version stored in the RCS file have not changed. However, most users would be happy to wait for an occasional unnecessary recompilation so long as they are happy that the correct version will be compiled.

A potentially more serious problem with the GNU MAKE rule is that it cannot tell which version is currently checked-out. For example, if the *COFLAGS* macro used with the previous *make* was *-r2* the source file in the working directory will be the most recent version from release 2. However, if you now perform a new *make* with the *COFLAGS* macro set equal to *-r1* you would expect to use the most recent version of the source files from release 1. Unfortunately the MAKE utility has no way of knowing that the source file in the working directory is not correct. If the RCS file was not modified since the last compilation was performed, the MAKE utility will incorrectly decide that the source file is up-to-date, because it was modified more recently than the RCS file.

Since the MAKE utility is not very well integrated with the MAKE utility, it is necessary to use a two step process when attempting to build an old version of your executables:

- Ensure that the correct versions of the source files are checked-out of RCS.
- Build the executables from these sources by using the MAKE utility. (This is the job for which MAKE was originally intended.)

A major advantage of using the MAKE utility is, that it will only perform the minimum number of re-compilation steps required to build the executables. However, if you blindly check-out all of the source files from RCS immediately prior to each re-compilation, the MAKE utility will be forced to redo all of the compilation steps. In order to avoid this problem, you must ensure that you only check-out source files when the currently checked-out version is incorrect.

You might ask: how can I ensure that ? You must follow the following steps:

- Check-out a copy of the desired version into a temporary file.
- Compare the contents of the temporary file with the current contents of the source file. If the files are different, replace the source file with the temporary file. Otherwise delete the temporary file.

The following shell script will automatically carry out the procedure for you.

```
#!/bin/sh
#
#   A shell script that only checks-out files
#   if the resultant version would be different from
#   the currently checked-out version
#
#   Brian O'Donovan 13-Apr-91
#
# define ${CO} to point to where the "co" command is on your system
# default is "/usr/new/co"
CO=/usr/new/co
#
#   First check the option flags specified
for i
do
    case $i in
        -*) OPT="${OPT} $1"
            shift;;
        *) break;
            shift;;
    esac
esac
```

```

done
echo "Performing a test check-out with options: ${OPT}"
#
#   The remaining parameters are files to be processed
while test $1
do
    echo "Processing $1"
    # check-out the correct version into a temporary file
    ${CO} ${OPT} -p $1 > /tmp/co$$
    if [ $? -ne 0 ]
    then
        echo "Error encountered while checking-out $1"
        exit 1
    fi
    # compare its contents with the source file
    diff $1 /tmp/co$$ > /dev/null 2> /dev/null
    # if they differ replace the source file
    if [ $? -ne 1 ]
    then
        echo "The source file $1 is being updated"
        mv /tmp/co$$ $1
    else
        echo "The correct version of $1 is currently checked-out"
        rm /tmp/co$$
    fi
    shift
done
echo "Processing Complete"

```

If you copy this script into a file named *test\_co*. You can then use the following commands to build release two of the executable file *prog*.

```

% test_co -r2 *.c *.h
% make prog

```

The main advantage of using this shell script is that the modification time of the file is not changed if the current contents are OK. I have designed an enhancement to the *co* command which would have the same effect; it is possible that this enhancement will be included in future release of RCS.

## 7

## RCS for Team Development Efforts

"Our military forces act as a team--in the game regardless of who carries the ball."

*Omar Bradley*  
1949

If you only use RCS to work on small, personal projects, the information from the last chapter should satisfy you. You know how to use RCS to preserve old versions, track the difference between successive versions, and so on. For a small project, this may be all you need to do.

However, RCS really comes into its own when you are managing a large project with several teams of developers. This common scenario places a premium on communications: keeping all sources in a common, central location, providing "interlocks" to prevent simultaneous updates, allowing check-ins to be coordinated. Unfortunately, this is one aspect of RCS usage that its standard documentation doesn't describe. And, to be completely honest, RCS does not solve the problem completely. Many users have developed scripts to provide the additional help that RCS needs, and several complete systems have been developed to enhance RCS's basic services.

In this chapter, we'll discuss what you need to do to manage large projects. We'll give you some help writing the kinds of scripts that are necessary for larger projects by showing you some sample scripts we have written. We will also discuss how to use some features of RCS itself (like states and version names) to make management of large projects more convenient.

You should be aware, however, some of the features which you might need to manage a large project are too complex for a simple shell script. Hence many systems have been written that have been layered "on top" of RCS. We will give an outline of a few such systems in the next chapter.

## Getting Everyone to Use the Same RCS Files

If you want to ensure that all of the project members use RCS library files from the same project's shared directory, you should get the system to automatically look in this directory for RCS library files. However this can cause problems in a large project, because each developer will have their own working files at different directory locations and RCS assumes that the library file is in a sub-directory named `./RCS` relative to the directory containing the working file.

If your operating system supports symbolic links, this problem is simple to solve. You simply create a symbolic link from each working directory to the project's RCS directory. For example the following sequence of commands will create symbolic links between the project's RCS directory (`/proj/lib/RCS`) and the directory locations `/usr/users/tom/RCS`, `/usr/users/dick/RCS` and `/usr/users/harry/RCS`.

```
% ln -s /proj/lib/RCS /usr/users/tom/RCS
% ln -s /proj/lib/RCS /usr/users/dick/RCS
% ln -s /proj/lib/RCS /usr/users/harry/RCS
```

After these symbolic links have been created, the users *tom*, *dick* and *harry* can access the project's RCS files as if they are in the `./RCS` sub-directory of their normal login directories (because the operating system considers these directory paths to be identical).

If your system does not support symbolic links, the problem of getting everyone to use the same RCS files is not quite as easy to solve. The only practical solution is for you to write a shell script which will intercept all of the RCS commands issued by the users and translate them into RCS commands that access the project's RCS directory. The following program shows you an example of a shell script which I wrote to accomplish this task. †

```
1  #! /bin/sh
2
3  comdir=/usr/new
4
5  RCSDIR='cat .rcsdir 2> /dev/null'
6  if test -z "$RCSDIR"
7  then
8      echo "The RCS directory has not been defined"
9      if test -d ./RCS
10     then
11         echo "It defaults to ./RCS"
12         RCSDIR="./RCS"
13     else
14         echo "Enter the name of the directory containing the RCS files"
15         read RCSDIR
16         echo $RCSDIR > .rcsdir
17     fi
18 fi
19
20 command=' echo $0 | sed 's/.*///g' '
21
22 for param in $*
23 do
24     case $param in
25
```

† The line numbers have been added for clarity and are not part of the script.

```

26         -*)
27             options="$options $param"
28             ;;
29
30         *)
31             libfile=$RCSDIR/` echo $param | sed 's/.*/g' ` ,v
32             files="$files $param $libfile"
33             ;;
34
35     esac
36 done
37
38 $comdir/$command $options $files
39 exit $?

```

The important part of this script is contained in lines 5-18. The script reads a file named *.rcsdir* to find out the name of the directory where the RCS files are contained. If no file named *.rcsdir* exists, the script looks to see if there is a sub-directory named *./RCS* available to store the RCS files (so that some users can continue to use RCS for their own personal projects). If no *./RCS* directory can be found, the user is prompted to type in the name of the RCS directory; this name is recorded in the file *.rcsdir* for future use.

The RCS directory name, which was deduced in lines 15-18, is used in lines 22-36. All command line parameters which do not start with '-' are assumed to be the name of a working file. For each working file, the name of the associated RCS library file is generated by pre-pending the name of the RCS directory and by appending ",v". The name of both the working file and the associated library file are given as parameters to the actual RCS command.

### NOTE

This shell script assumes that you will only give the name of the working file when you call RCS commands and not the name of the library file. Strange things will probably happen if you specify the name of the library file.

The smart part of this shell script is on line 20 where it figures out the name of the RCS command to use from the command used to invoke the shell script. This means that the same shell script can be used as a pre-processor for all of the RCS commands. This script should be installed as */usr/bin* so that it will be seen earlier in users' search paths than the real RCS commands which will be in */usr/new*. There should be links to this script from each of the RCS command names. The following commands show how you should install this script. (You will probably need to be logged in as *root* before you are allowed to execute these commands.)

```

% cp SCRIPT /usr/bin/rcs
% cd /usr/bin
% ln rcs co
% ln rcs ci
% ln rcs rlog
% ln rcs rcs-clean
% ln rcs rcs-merge
% ln rcs rcs-diff
% chown root.system rcs co ci rlog rcs-clean rcs-merge rcs-diff
% chmod 0755 rcs co ci rlog rcs-clean rcs-merge rcs-diff

```

## Getting Everyone to Use the Same Working Files

We said earlier that whenever someone wants to change one of the projects' source files, they should:

- Check-out their own working copy of the file.
- Modify the source file.
- Test the modified source file.
- Re-modify and re-test the source file until they are happy with the results of the tests.
- When they are happy with the test results, they can check-in the modified source file to the shared directory.

However, in most software systems it is not possible to test a single source file on its own. Hence, a user must check-out a read-only copy of all of the source files in addition to a writable copy of the one he/she is actually going to change. This situation is unsatisfactory because it is wasteful on disk space and because it means that some users could easily be using old versions of the source files if the files are updated while after they check-out a read-only copy.

It would be preferable if all of the project's members could share copies of the files that they are not changing as well as having their personal copies of the files that they intend to modify. This set-up is possible by using symbolic links and two simple shell scripts to maintain them. The first shell script (which is shown below) is called *make\_links*. It will establish symbolic links between your current working directory and each of the source files the project's shared directory.

```

1  #! /bin/sh
2
3  RCSDIR=`cat .rcsdir 2> /dev/null`
4  if test -z "$RCSDIR"
5  then
6      echo "RCSDIR has not been defined"
7      if test -d ./RCS
8      then
9          echo "It defaults to ./RCS"
10         RCSDIR="./RCS"
11     else
12         echo "Enter the name of the directory containing the RCS files"
13         read RCSDIR
14         echo $RCSDIR > .rcsdir
15     fi
16 fi
17
18 case $RCSDIR in
19
20     */RCS)
21         rcsdir=$RCSDIR
22         pubdir=` echo $RCSDIR | sed 's//RCS$//' `
23         ;;
24
25     *)
26         rcsdir=${RCSDIR}/RCS
27         pubdir=$RCSDIR
28         ;;
29
30 esac
31
32 count=`ls -la $pubdir 2> /dev/null | wc -l`

```

```

33  if test $count -le 3
34  then
35      echo "There are no files in
36      exit 1
37  fi
38
39  for file in $pubdir/*
40  do
41      ln -s $file .
42  done
43
44  exit 0

```

The *make\_links* script uses the same logic as the script shown in the last section to locate the project's shared directories. It also assumes that the project's working files are stored in a directory which is one level above the directory containing the project's RCS library files.

The *make\_links* script should be run once by each user in the directory in which he/she intends to do development work. For example, if the project had three source files *prog.c*, *prog.h* and *utils.c*. The project's shared directory would contain three RCS library files */proj/lib/RCS/prog.c,v*, */proj/lib/RCS/prog.h,v* and */proj/lib/RCS/utils.c,v*. The project would also have three publicly accessible working files */proj/lib/prog.c*, */proj/lib/prog.h* and */proj/lib/utils.c*. After *fred* runs the *make\_links* script in his home directory */usr/users/fred*, he can then access the three publicly accessible working files */proj/lib/prog.c*, */proj/lib/prog.h* and */proj/lib/utils.c* as if they were named */usr/users/fred/prog.c*, */usr/users/fred/prog.h* and */usr/users/fred/utils.c*.

These symbolic links are fine for files that *fred* does not intend to modify, but he will need real copies of any files that he intends to update. What he needs is a script which will automatically delete the symbolic link to the shared copy of the working file immediately before he checks-out his own copy. This script should also automatically re-create the symbolic link and check-out a shared copy of the file immediately after he checks-in a new version.

The */usr/bin/rcs* pre-processor script which we described in the last section can easily be modified to do this extra work. The following is an example of how this can be done:

```

1  #! /bin/sh
2
3  comdir=/usr/new
4
5  RCSDIR=`cat .rcsdir 2> /dev/null`
6  if test -z "$RCSDIR"
7  then
8      echo "RCSDIR has not been defined"
9      if test -d ./RCS
10     then
11         echo "It defaults to ./RCS"
12         RCSDIR="./RCS"
13     else
14         echo "Enter the name of the directory containing the RCS files"
15         read RCSDIR
16         echo $RCSDIR > .rcsdir
17     fi
18 fi
19
20 case $RCSDIR in
21     */RCS)
22

```

```

23         rcsdir=$RCSDIR
24         pubdir=` echo $RCSDIR | sed 's//RCS$//' `
25         ;;
26
27     *)
28         rcsdir=${RCSDIR}/RCS
29         pubdir=$RCSDIR
30         ;;
31
32 esac
33
34 command=` echo $0 | sed 's/.*/g' `
35
36 for param in $*
37 do
38     case $param in
39
40         -*)
41             options="$options $param"
42             ;;
43
44         *)
45             libfile=$rcsdir/` echo $param | sed 's/.*/g' `,v
46             files="$files $param $libfile"
47             local_files="$local_files $param"
48             ;;
49
50     esac
51 done
52
53 case $command in
54
55     co)
56         rm -f $local_files ' 1> /dev/null 2> /dev/null'
57         $comdir/co $options $files
58         exit_status=$?
59         ;;
60
61     ci)
62         $comdir/$command $options $files
63         exit_status=$?
64         for file in $local_files
65         do
66             pubfile=$pubdir/` echo $file | sed 's/.*/g' `
67             rcsfile=$rcsdir/` echo $file | sed 's/.*/g' `,v
68             $comdir/co $pubfile $rcsfile
69             ln -s $pubfile $file
70         done
71         ;;
72
73     *)
74         $comdir/$command $options $files
75         exit_status=$?
76         ;;
77 esac
78
79 exit $exit_status

```



The main changes to this script are concentrated around lines 53-70 where the *co* and *ci* commands are singled out for special treatment. Line 56 deletes the symbolic link immediately before the actual *co* command is executed. Lines 66-68 check-out a new working file for the project's shared directory and line 69 re-creates the link from the current directory to it.

There is one major problem with sharing working files by means of symbolic links as we describe in this section; the *make\_links* procedure assumes that you always want to use the default version of each working file. It is not possible to use this working arrangement if you want to work on a branch other than the default.

Another problem is that it may be difficult for you to track down bugs in one module of the software while other project members are simultaneously making changes to related code modules. Therefore, if you are concerned to ensure that nobody else modifies any of the modules while you are testing, you should check-out your own personal copy of all files.

### What is Missing ?

The shell scripts described in this chapter, can greatly enhance the usefulness of RCS for team based projects. In many ways these scripts achieve much of the functionality of a system such as CCSLAND (which is described in the next chapter). However, these shell scripts are still missing many features that you might want. The two most significant shortcomings of these shell scripts are:

- These scripts do not allow you to manipulate entire directory hierarchies with one command.
- These scripts do not facilitate development teams where some of the developers are working on different variants of the code from the others.

However, it is not really practical to implement anything more complex in a simple shell script. If you need this enhanced level of functionality, you should consider getting a system such as CCSLAND.

# 8

## Securing Your RCS Files

"The prince who relies upon their words, without having otherwise provided for his security is ruined; ..."

*Niccolo Machiavelli*  
1532

If you use RCS for your own personal project, you can solve all potential security problems by setting the security on the RCS library file such that only you can use it. However, this simple solution will not work for large projects because all of the project members must be able to access the project's files. On a large project you will probably be concerned to ensure that:

- Only project members can access the source files.
- Only certain named project team members are allowed to update some of the source files.
- All interaction with the project's RCS files uses the approved processes.

RCS provides you with a mechanism called *access control lists* which allow you to restrict who may update your RCS files. However, the RCS access control lists are not a totally effective way of securing your RCS files. In order to be really effective, RCS access control lists must be used in conjunction with normal UNIX files protection. In this chapter we will show you how to use RCS access control lists; we will then describe how to use normal UNIX file protection to further enhance the security of your RCS files.

## Access Control Lists

An RCS file may contain an *access control list*. This list names the users who are permitted to update the file. When a RCS file is first created, it contains no access control list. By definition, a null access control list means that anyone can update the RCS file, subject to normal UNIX file protection.

Users can be added to the access control list of a file using the *rcs -a* command. For example the following command creates an access control list for the RCS file *RCS/sample.c,v* containing the users *mary* and *jane*.

```
% rcs -amary,jane sample.c
```

You can also use the *rcs -a* command to add users to an existing access control list. For example the next command, adds the users *fred* and *harry* to the existing access control list for the RCS file *RCS/sample.c,v*. It now contains the users *mary*, *jane*, *fred* and *harry*.

```
% rcs -afred,harry sample.c
```

Once users have been added to the access control list, they may be deleted with the *rcs -e* command. For example the next command deletes the users *fred* and *mary* from the access control list for the RCS file *RCS/sample.c,v*. It now contains the users *jane* and *fred*.

```
% rcs -efred,mary sample.c
```

If you use the *rcs -e* command without giving any user names you will delete the entire access control list for the file. For example, the following command deletes the entire access control list for the RCS file *RCS/sample.c,v*. Now any user can access it, subject to normal UNIX file protection.

```
% rcs -e sample.c
```

You are likely to want to have the same access control list on more than one file. Assume that your project has 200 source code files. You might want to give some of the members of your development team access to any file, but you want to prevent others from accessing the files. In this case, the access control list for each file would be the same. To save you from giving a lengthy *rcs -a* command 200 times, RCS lets you set up the access control list for one file, and then copy that list to the others.

The *rcs -A* command performs the copying. For example, either of the following two commands will copy the access control list of the file *RCS/other.c,v* into the file *RCS/sample.c,v*

```
% rcs -ARCS/other.c,v RCS/sample.c,v
% rcs -Aother.c sample.c
```

If you give a list of file names to the *rcs -A* command, it will copy the access control list to all of the named files in one go. The following two commands show how you might use this feature.

```
% rcs -Aother.c sample.c test.c sample.h
% rcs -ARCS/other.c,v /pub/RCS/*,*v
```

Access control lists are a very powerful and flexible mechanism for restricting update access to your RCS files, because you can explicitly name each user who is allowed to update each of the RCS files. However, there are a few key points that you should remember about the use of access control lists.

- RCS access control lists only restrict update operations, they do not place any limit on who may read the file. Anyone with UNIX read permissions for the RCS file can read the file.
- RCS access control lists act in conjunction with the UNIX file access permissions. In order to update a RCS file, it is necessary to have write access to the directory containing the RCS file **and** be on the access control list (if one exists). If a user does not have write permission to the directory containing the RCS file, they will not be able to update the file, even if he/she is on the access control list.
- Regardless of the access control lists, the super-user ("root") and the file's owner are always allowed to update a file. These two users are exempt from any restrictions imposed by an access control list. In other words: you cannot protect a file against a malicious user with the superuser password, and you cannot protect a file against its owner. You should also remember that, the access control list only restricts who can update the file; it has no effect on who can read the file.
- RCS access control lists are really only suitable for restricting access in a friendly environment. It is relatively simple to circumvent the RCS access control mechanisms. However, it should not be possible to circumvent the UNIX file access protections unless there is a bug in the operating system. Therefore you should use UNIX file access mechanisms whenever you are really concerned about security.

## Using UNIX file protection

Another mechanism for securing access to your RCS files, is the file protection mechanism of the operating system itself. In fact, the only way of restricting read access to your RCS files is by setting the UNIX file permissions on the RCS files. For example, if you would like to ensure that only you can read your RCS files, you should use the following command to restrict UNIX read access to the files.

```
% chmod 0400 RCS/*
```

Likewise, you can restrict read access to your RCS files to users that are in your group with the following command.

```
% chmod 0440 RCS/*
```

The default access permission for a newly created RCS file, is taken from the access permissions of the working file. The RCS file will have the same access permissions as the working file except that all write and execute permission is turned off. Therefore, the default rule is that anyone who could read the working file will be able to also read the corresponding RCS file. Nobody is allowed to directly write or execute a RCS file, because it would make no sense for them to do so.

RCS has a special way for dealing with updates to library files. Whenever you use a RCS command to update a library file (*e.g.*, *co -l*, *ci* etc.) RCS does not directly update the existing file. Instead, RCS initially creates a temporary file containing your update; when the command is ready to complete, the temporary file is used to replace the original file. This strategy is adopted so that, if the system crashes, you will never be left with a RCS library file where an update has been only partly carried out. However, there are a number of significant side effects of this strategy which you must keep in mind:

- Since a new RCS file is created every time, the owner of the RCS file is the person who last updated the file. This means, for example, that the last person to update a file can change its UNIX access permissions.
- UNIX write permissions on a RCS file are redundant, because the system never directly writes into RCS files. This is the reason why RCS automatically turns off the write permission on all RCS files. (Turning off write permission has no effect upon the operation of the RCS commands, but it discourages users from directly editing the RCS library file.)
- If you have a hard link to a RCS file, your link will always point to the same copy of the RCS file. This is because hard links are associated with the actual i-node number of the file and every time a RCS file is updated it gets a new i-node number (because it is actually a new file).

If you want to update a RCS file, you must have write access to the directory containing the RCS file (in order to be able to create the replacement file). You can restrict update access to your RCS files by setting the appropriate UNIX protection on the directory containing your RCS files (not the by setting the protection on the RCS library files themselves). For example, the following command will only allow users in the same group as yourself to update the RCS files.

```
% chmod 0660 RCS
```

Users can be denied permission to update an RCS file, either because they don't have write access to the directory, or because they are not on the access control list. Therefore the previous command has the effect of restricting update access to users who are in the same group as yourself **and** are on the access control list.

## Using the Set-uid Bit

All of the RCS commands obey the restrictions of the access control list. However, malicious users can easily circumvent the access control mechanism by using some non-RCS commands to manipulate the RCS library files. It would be a good idea if you could ensure that users only used the legitimate RCS commands to manipulate the RCS files. At first glance this seems like an impossible dream; but by using the set-uid bit you can sometimes make this dream a reality.

In most variants of the UNIX operating system, files have a set-uid protection bit. When this bit is turned on, anyone executing the program contained in this file, will have the rights and privileges of the owner of the file. This mechanism is useful when you want to allow users to perform some specific update operations that they would not normally be allowed to carry out.

If you wish to use the set-uid bit with the RCS commands, you must carry out the following steps:

- Create a special account for use with RCS (*e.g.*, a user named *rsc\_user*). It is normally a good idea to disable logins on this account.
- Set the ownership and protection on all RCS directories so that only the new account has write permission for the directory. *e.g.*:

```
% chown rsc_user ./RCS /proj/RCS /alt/RCS ...
% chmod 0644 ./RCS /proj/RCS /alt/RCS ...
```

- This account should be given ownership of the files containing the RCS commands. *e.g.*:

```
% cd /usr/new
% chown rcs_user co ci rcs ....
```

- The set-uid bit should be set on the RCS program files. (You may need to be the super-user before you are allowed to do this.) *e.g.*:

```
% chmod u+s,-w co ci rcs ....
```

### CAUTION

It is important that users should not be given write permission to a file which has its set-uid bit set. This is because, if write permissions enabled on such a file, user can overwrite the program with any program they wish. They will then be able to execute this new program with the enhanced permissions.

When the RCS programs have been configured to run with the set-uid bit, all RCS files will be owned by the *rcs\_user* account. You can ensure that users can only update the RCS files through the standard RCS commands, by changing the ownership and access permissions on the directories containing the RCS files. For example:

```
% chown rcs_user RCS
% chmod 0700 RCS
```

When file protection has been configured this way, users cannot hack at the RCS files because they have no access permission for the RCS files themselves. However, users can update the RCS files by using the regular RCS commands because the set-uid bit gives them the necessary extra permissions. (This configuration also ensures that the RCS access control lists are respected).

Unfortunately, RCS does not always work correctly with the set-uid bit set. Problems can be encountered due to bugs in either the RCS code or the underlying operating system. The following simple tests will let you know whether or not the set-uid feature is working correctly on your system.

- Create a new RCS file and then check that the file belongs to the *rcs\_user* account. For example:

```
% rcs -i RCS/test,v
% ls -l RCS/test,v
```

- Set-up the protection on the RCS directory so that you do not have write access to the directory. Attempt to update a RCS file: you should succeed due to the fact that set-uid bit gives you extra permission.
- Check-out a RCS file with a lock. Look to see that the lock has been granted in your name (and not the name of the *rcs\_user* account). Also look to see that the resultant working file belongs to you (and does not belong to the *rcs\_user* account).

If any of these simple tests fail, you should not run RCS with the set-uid bit set. The latest version of RCS (5.6) allegedly can be used on any variant of UNIX with the set-uid bit turned-on, you should try to get this version if your current version will not work.

## Using Groups to Restrict Access

One of the main dis-advantages to using the set-uid bit is that everyone is treated equally from the point of view of UNIX file permissions. This is not a problem for update access because the access control list can be used to limit who has update access to the RCS files. However, this can cause you to unwittingly give read access to everyone.

If you are using the set-uid bit on the RCS commands, the owner of the RCS file will be *rsc\_user*. If owner read is enabled on your RCS files, anyone will be able to read your files with the *co* command (because the set-uid bit ensures that anyone who executes the *co* command is treated as if they were *rsc\_user*).

It is normal for most UNIX file to allow reads by the owner of the file. However, if a RCS library file is created from such a file, the default will be for the RCS library file to have owner read enabled. If you do not want to let everyone have read access to your RCS files, you must ensure that owner read permission is turned off on all of your RCS files.

If you are using the */usr/bin/rsc* pre-processor script described in the previous chapter, it is quite simple to add a command to automatically turn off owner read permission to RCS files. The following command should be added immediately after the *ci* command (line 62).

```
% chmod o-r,u-r $rscfile
```

If owner read is disabled on your RCS files, it is the group access permissions which decide whether or not you are allowed to access the RCS files. Read access will be limited to members of the same UNIX group as the person who created the file. Hence, if you want to limit read access to project team members, you should put all of the project team members into their own group.

# 9

## RCS Based Systems

"Obviously, a man's judgment cannot be better than the information upon which he has based it."

*Arthur Hays Salzberger*  
1948

RCS was developed to store and manage multiple versions of a single file. Some recent additions to the RCS system, (such as the *rcsclean* command), allow the RCS system to efficiently manage multiple files, but only if all of the files are in the one directory. Many Large software software projects involve a complex directory hierarchy of source files: unfortunately RCS is not very good at managing these directory hierarchies.

In addition, the exclusive locking system employed by RCS works on the assumption that only one developer should be modifying a source file at the same time. However, in large projects it is often not practical to abide by this restriction. RCS is not good at co-ordinating the situation where multiple developers are simultaneously modifying the same source files.

Because of these deficiencies in RCS, many large projects feel that they need tools with greater functionality. Rather than develop these new tools from scratch, many developers have adopted the strategy of providing this additional functionality in a software package that runs on top of RCS itself. Such tools provide for effective management of multiple simultaneous updates to a complex directory hierarchy of source files. However, they call upon RCS itself for the basic version management functions.



In this chapter we will review three such tools which are implemented as an additional layer of software on top of RCS itself. The Software Project Management System (SPMS), the Configuration Management System (CCSLAND) and the Concurrent Versions System (CVS). We will also introduce you to a system called Distributed RCS (DRCS) which is an enhanced form of RCS for use by distributed development teams.

## Software Project Management System (SPMS)

We mentioned earlier that, RCS is capable of managing multiple files in one directory, but it has problems dealing with multiple directories. RCS is not alone in this regard; many UNIX utilities (such as MAKE) also have difficulty in dealing with files in more than one directory. SPMS can be used with any existing tool to allow the tool to work with multiple directories as easily as one directory.<sup>†</sup>

SPMS works by grouping several directories into one single project. For each project, SPMS maintains a special "project database" in the project's root directory. The project database contains a list of all the directories that are part of the project. SPMS has a number of commands which act on the project directories. The simplest of these is the *pexec* command.

The *pexec* command will execute any arbitrary command in each of the directories in turn. For example, the following command will execute the command **rcsclean \*.c** in each of the project directories (effectively doing a clean up of all the project sources).<sup>‡</sup>

```
% pexec 'rcsclean *.c'
```

Likewise you could check-out release 2 of all the project's source file by using the command:

```
% pexec 'co -r2 *'
```

If you were not using SPMS you would have to *cd* into each of the project's directories and issue the *co* command from each individual directory. *e.g.*

```
% cd src
% co -r2 *
% cd ..
% cd docs
% co -r2 *
% cd ..
% cd man
% co -r2 *
% cd ..
etc.
etc.
```

<sup>†</sup> It is easy to get confused by the name *Software Project Management System*. In the SPMS terminology, a project is just a directory hierarchy. SPMS is not capable of performing the tasks that you would normally associate with a project management system such as, creating schedules or tracking the progress of the development work.

<sup>‡</sup> You will notice that the command **rcsclean \*.c** is placed within quotation marks when used with the **pexec** command. This is to stop the shell from expanding the *\*.c* term into all of the files in your current working directory whose names end in *.c*.

As well as storing a simple list of all the project directories, SPMS also stores a type and priority for each directory. The directory type can be used to select which directories the command is to use. For example, it is only necessary to execute the *make* command in project directories which contain source code. Hence you could use the following the following command will perform a *make* command in all of the project's directories whose type is *src*.

```
% pexec -Tsrc make
```

The priority is used to specify in which order the directories should be processed. Directories with a lower priority are guaranteed to be processed first. This can be important if the *make* process in one directory requires to use the output of the *make* process in another directory.

You can create a project database by using the *mkproject* command. SPMS maintains a project hierarchy, which is analogous to the UNIX directory hierarchy. The *chproject* command is provided to activate a project (*i.e.*, make that project the current working project; *chproject* performs the SPMS equivalent to the UNIX *cd*). The project database created by the *mkproject* command initially contains no directories. Directories can be added to the project database with the *pmkdir* command. For example the following series of commands creates a new project named *util*, activates that project and then adds a new directory to the project with a name *man*, a type *doc* and a priority 2.

```
% mkproject -d util
% chproject util
% pmkdir -Tdoc.2 man
```

SPMS has its own semantics for representing project hierarchies which mirrors the UNIX semantics for representing the directory hierarchy. The symbol *^* is used to represent the root of your project hierarchy; it is analogous to */* character which represents the root of the directory hierarchy. For example the project name *newproj* represents a project which is a child of the currently active project. However, a project name *^newproj* represents a project which is a child of your root project.

Similarly, the *^* character can be used as a separator in project names. Thus a project name *main^part1* refers to a project named *part1* which is a sub-project of the project named *main* which is a sub-project of your current working project.

By convention, the UNIX file named *.* refers to your current working directory and the file named *..* refers to the directory one above it in the directory hierarchy. Not to be outdone, SPMS has introduced the convention that the name *...* can be used to refer to the currently active project and the name *....* refers to the project immediately above it in the SPMS project hierarchy. (If you are not confused now you never will be.)

As well as the general purpose *pexec* command. SPMS also has some special purpose commands for performing commonly used operations. These commands are summarized in the following table.

Table 9-1: SPMS specialized commands

Command	Function
<b>pdiff</b>	Similar to the UNIX <i>diff</i> command. It compares the contents of the files in two projects (e.g. the command <b>pdiff ^proj1 ^proj2</b> compares each of the files in the project named <i>^proj1</i> with the corresponding file in the project named <i>^proj2</i> )
<b>pgrep</b>	Similar to the UNIX <i>grep</i> command. It searches for a pattern in files in any of the project's directories. (e.g. the command <b>pgrep install makefile</b> searches for occurrences of the pattern <i>install</i> in files named <i>makefile</i> in any of the current project's directories).
<b>pfind</b>	Similar to the UNIX <i>find</i> command. It finds files in a project. (e.g. the command <b>pfind prog.c</b> will print out the location of any file named <i>prog.c</i> in any of the current project's directories).

It is clear that the benefit of using the SPMS tool increases directly in line with the number of directories in the project. If your file are contained in a small number of directories, you will not derive any benefit from using SPMS. However, if you are working with a project which has a complex directory hierarchy (such as the sources for UNIX itself), you will probably find SPMS very useful.

I will not discuss SPMS in any greater detail here because there is ample documentation provided with the system itself. The source code of SPMS is available free of the normal licensing restrictions. SPMS was included with release 4.2 of BSD UNIX. However, at the time of writing no decision had been made about how SPMS would be distributed separately from BSD UNIX. If you contact the author Peter Nicklin at the address below, he will tell you how you can get a copy of SPMS.

Peter Nicklin,  
Hewlett-Packard Company,  
Circuit Technology Research and Development,  
Circuit Technology Group,  
5301 Stevens Creek Boulevard,  
PO Box 58059,  
Santa Clara,  
California 95052-8059,  
USA.

Phone: (408) 246 4300  
Fax: (408) 249 6255  
Email: [nicklin%hpdtc@hplabs.hp.com](mailto:nicklin%hpdtc@hplabs.hp.com)

## CCSLAND

CCSLAND is a system which allows you to perform version management tasks on an entire directory hierarchy rather than just on a single source file. CCSLAND uses the services of RCS for version control of the individual files. CCSLAND extends the functionality of RCS by grouping a complex directory hierarchy into a single *source repository* which can be managed as a single entity. The letters CCS at the start of CCSLAND stand for **C**onfiguration **C**ontrol **S**ystem.

You already learned that SPMS can be used to manage a complex directory hierarchy as a single entity. However, CCSLAND has many additional functions to aid in the management of a complex set of source files. The most significant functionality is, to allow each user to establish their own workspace containing a virtual copy of the source repository. This functionality is very useful because it allows you to experiment by making changes in your own workspace before making these changes publicly available in the source repository.

I say that CCSLAND creates a virtual copy, because CCSLAND actually establishes symbolic links from your workspace to the files in the project's public directories, rather than actually copying the files. This saves a substantial amount of disk space relative to what would be required to actually copy the entire directory hierarchy of the project. In addition, symbolic links have the advantage that you will automatically see the effects of any updates to the project's public files. Naturally, CCSLAND enforces the rule that you must create a copy of any file that you intend to modify.

If you wish to work on files in a project which uses CCSLAND, the first step is to create your own private workspace containing a cloned copy of the relevant portion of the project hierarchy. The CCSLAND command to do this is *ckclone*. For example, the following command will create for you a private workspace containing a cloned copy of all the project files in the project's *utils/uucp* directory (and all sub-directories of it).

```
% ckclone utils/uucp
```

The cloned directories are placed at a directory path named *utils/uucp* relative to your current working directory. For example, if your current working directory is *~/proj/workdir* the cloned directories will be placed at *~/proj/workdir/utils/uucp*.

You might well have asked: how did CCSLAND know where the project's public directories are located? The answer is simple: it reads this information from a special CCSLAND configuration file named *.CCS*. The *.CSS* file contains information about the location of the public directory as well as selecting certain options which will control how the CCSLAND commands operate.

The most common type of line in a *.CCS* file contains something like:

```
cklockflags=-f
```

This line tells CCSland that by default you would always like to use the *-f* option with the *cklock* command.

You can have several *.CSS* configuration files. The CCSLAND system will first look for a *.CSS* file in the root (*/*) directory. It will then look for *.CSS* files in each of the components of your current pathname. If options specified in later *.CSS* files conflict with options specified in earlier *.CSS* files, the *.CSS* file which was read last is the one which has effect. The way that *.CSS* files are processed in sequence means that the *./CSS* file specifies system wide defaults and a *.CCS* file in a users home directory specifies the defaults for that user. A user can activate a different set of defaults by changing their working directory to a sub-directory containing new defaults.

Once you have created your private workspace it is necessary to check-out real copies of the files you intend editing. It is also necessary that you should place a RCS lock on them. You could use the RCS *co -l* command to do this, however, this would require you to know the full pathname of the RCS files in the project library. In order to make this easier for you, CCSLAND provides the *cko* command which is a special interface to the RCS *co* command. The main difference between the *co* and *cko* commands is that, while *co* command will look for the RCS file in the *./RCS* directory, *cko* will look for the RCS file in the project directory determined by reading the *.CCS* file(s). In addition, the *cko* command will delete the symbolic link between your directory and the associated file in the project directory which was established by the *ckclone* command.

When you are finished editing the files, you must check-in the updated versions into the project's RCS files. Once again you could use the standard RCS *ci* command to do this. However, it is more convenient to use the CCSLAND *cki* command. Like the *cko* command, the *cki* command will look for the RCS file in the relevant project directory. In addition the *cki* command will re-establish the symbolic link between your workspace and the public copy once your local copy has been deleted by the *ci* command.

CCSLAND has a number of other commands which act as a front end to the RCS system. The following table summarizes the functionality of each of these and shows which RCS command contains approximately the same functionality.

Table 9-2: CCSLAND interface to RCS

CVS Command	RCS Equivalent	Function
<b>cko</b>	<b>co</b>	Check-in a file
<b>cki</b>	<b>ci</b>	Check-out a file
<b>ckdiff</b>	<b>rcsdiff</b>	Displays the differences between two versions of a file
<b>cklock</b>	<b>rcs -l</b>	Locks a file ( <b>cklock</b> can also lock an entire directory hierarchy).
<b>cklog</b>	<b>rlog</b>	Shows the modification history of a file

CCSLAND recognizes that, as well as having multiple developers working on the same project, it is also possible that developers can be working on different incarnations of a project at the same time. In order to support this, CCSLAND has the concepts of *checkpoints* and *configurations*. Each source file will only have one RCS file, however, it may have several checkpoints and each of the checkpoints may have several configurations.

A checkpoint consists of a particular variant of the software sources. For example you might have a *terminal* checkpoint (which is the variant of the source files which work on regular terminals) and a *X\_win* checkpoint (which is the variant of the source files which work on regular terminals). There will be a branch created in the RCS files for each checkpoint.

A configuration consists of the compiled version of the source files, compiled for a specific hardware/software environment. For example you might have a *VAX* configuration (your executables compiled to run on VAX hardware) and a *Sun* configuration (your executables compiled to run on Sun hardware).

You can specify in your *.CCS* file which configuration and checkpoint you are working on. For example, if you are working on the *VAX* configuration of the *X\_win* checkpoint, you are using the *X\_win* branch of the source files and you are compiling it to work on VAX systems. The CCSLAND system will ensure that when you make changes to the sources, these changes are not seen by developers who are working on a different checkpoint. Likewise, if you compile a copy of the executable files, this will only be seen by the users who are working on the same configuration of the same checkpoint.

In addition to the management of private workspaces, CCSLAND also contains a utility called *makenv* which automatically generates a complete *Makefile* from a simple template file. This utility is very useful because it frees you from the tedious and error prone task of manually updating the *Makefile* every time the dependencies are changed. However I will not describe this utility here because it is functionality is not directly related to version control.

In summary, CCSLAND is a tool that can be layered on top of RCS to simplify the management of large complex projects. You will only really need to use a tool such as CCSLAND if your project contains a complex directory hierarchy and has multiple developers (otherwise RCS itself will be adequate). However, it is conceivable that you might want to get a copy of CCSLAND just to have the *makefile*

generator *makenv*.

Unlike both SPMS and CVS, CCSLAND is not available for free. The author describes it as *CHEAPWARE* because he only charges \$2K for a copy of the source code and a single site license. I will leave it to you to decide whether or not this really constitutes *CHEAPWARE*. You should contact Nathaniel Bronson at the address below if you are interested in either placing an order for CCSLAND or just getting more information about it.

Nathaniel Bronson,  
Microvation Consultants,  
20 Sperry Road,  
Madison,  
CT 06443,  
USA.

e-mail: tan@microvation.com  
fax: (203) 421 5292

## Concurrent Versions System (CVS)

The Concurrent Versions System (CVS) † is a software package that is layered on top of RCS. ‡ However, RCS and CVS have totally different philosophies with regard to how to prevent conflicting updates. While RCS enforces a strict rule that no two users may simultaneously update a source file, CVS takes a more pragmatic view that it is sometimes necessary to allow two users to simultaneously work on the same file. CVS attempts to repair conflicts after they happen, by using the *rcsmerge* command, rather than prevent them happening in the first place.

The basic principle of CVS is that all of the project's developers share a single source repository containing the RCS library files; however each developer has their own copy (version) of the working files. The name *Concurrent Versions System* arises from the fact that multiple developers can be concurrently modifying their own version (copy) of a source file. I prefer to say that each user has their own copy rather than use the CVS terminology that each user has their own versions. This is to avoid a potential confusion between RCS versions and CVS versions.

If you wish to modify certain files from the project's source repository, you must first establish your personal copy of the files with the **cv**s **checkout** command. It is normal to checkout a copy of an entire section of the project's directory hierarchy rather than just the files you wish to modify. This is so that you will be able to compile and test the modified files.

---

† In this section I am actually describing CVS II. The original CVS system (CVS I) was developed by Dick Grune as a set of shell scripts in 1986. All of the concepts implemented by Dick Grune in his original system were incorporated along with a number of enhancements by Brian Berliner and Jeff Polk when they developed CVS II in 1989. While the original CVS concepts are retained by CVS II, the actual commands and operations were changed substantially.

‡ CVS normally calls the RCS routines to manipulate the RCS library files on its behalf. However, CVS does sometimes directly manipulate the RCS library file. For this reason you must be careful that the version of CVS you are using is compatible with the version of RCS that you are using.

You must tell CVS where the project's source repository is located. You can do this either by defining an environment variable *CVSROOT* or by using the **-d** option with the **cvs** command. The following examples show the two ways of telling CVS that you would like to establish your personal copy of the *bin/cat* section of the project's source repository located at */src/proj*. The copies will be created in a directory named *bin/cat* relative to your current working directory, by performing a RCS *co* command on the library files located in the directory */src/proj/bin/cat*.

```
% cvs -d /src/proj checkout bin/cat
```

or

```
% setenv CVSROOT /src/proj
% cvs checkout bin/cat
```

In addition to creating the copies of the source files themselves, the *checkout* command will also create a special sub-directory *bin/cat/CVS.adm* to contain CVS related status information about your copies of the source files. The information contained in the files in this sub-directory is vital to the correct operation of subsequent CVS commands.

Once you have established your copies of the source files with the check-out command, you are free to edit them as you wish. However, the corresponding RCS library files are not locked and anyone else can be simultaneously modifying their copy of the same file. As a result, it is necessary for you merge your changes with the changes (if any) that the other developers have been simultaneously making, before you return the files to the source repository.

The **cvs update** command is provided to automatically update your copy of the source files in line with the changes that anyone else has made to the files in the source repository. There are three possibilities about what the *update* command will do:

- If no new versions were checked-in to the RCS library file since you performed the **cvs checkout** command: no action is taken.
- If a new version was checked-in to the RCS library file since you performed the **cvs checkout** command, but you made no modification of your own: the **cvs update** command will replace your copy of the file with the new version from the source repository.
- If a new version was checked-in to the RCS library file since you performed the **cvs checkout** command, and you made modifications of your own: the **cvs update** command will merge your changes with the changes which were applied to create the new version in the source repository. This is done with the *rcsmerge* command.

The success of the CVS policy of allowing simultaneous updates, depends upon the reliability of the *rcsmerge* command. In most real life situations, you will rarely need to call the *rcsmerge* command. Even when you do call the *rcsmerge* command, it will produce a sensible result in as many as 90% of all merges. However, it is inevitable that *rcsmerge* will sometimes produce unexpected results. When deciding to use CVS, you must balance the inconvenience caused by RCS locking against the inconvenience of the occasional failure of the *rcsmerge* based strategy. If you do decide to use CVS, you should be constantly vigilant for strange output from the *rcsmerge* command.

Once you have made your changes to the source files and used the **cvs update** command to incorporate other people's changes, the next step is to do a RCS *ci* to place the newly created versions into the RCS library files in the source repository. The CVS command to do this for you is *commit*. For example, the following command will check-in the modified files from your copy of the *bin/cat* sub-section of the project's hierarchy into the RCS files in the project's source repository.



```
% cvs commit bin/cat
```

The **cvs commit** command, will first check to see that your source files have incorporated everyone else's changes by means of the **cvs update** command, it will then check-in the modified files by calling the RCS *ci* command. The **cvs commit** command, will only check-in the files that you have actually modified. The RCS library files will only be locked immediately before the new version is checked-in.

The *checkout*, *update* and *commit* commands form the essential core of the CVS system. However, there are several other utility commands provided with CVS. I do not have room to describe all of the CVS utility commands here, but I will briefly describe the *tag* and *patch* commands because they are particularly useful.

The *tag* command can be used to assign a symbolic name to the appropriate version of all of the files used in a release. For example the following command causes the symbolic name *alpha\_release* to be assigned to the current default version of all files in the *src* project directory (and all sub-directories of the *src* project directory).

```
% cvs tag alpha_release src
```

If you later want to track down some bugs that were reported to occur in the *alpha\_release* of your software, you can use the following command to retrieve the correct version of all of the source files.

```
% cvs checkout -r alpha_release src
```

The *tag* command can also be used with the *patch* command to build patches for your software. For example, if you had two releases names *alpha\_release* and *beta\_release*. The following command would create a file named *alpha\_to\_beta* which can be used with Larry Wall's **patch** program to update a copy of the *alpha\_release* version of your software to a copy of the *beta\_release* version.

```
% cvs patch -r alpha_release -r beta_release src > alpha_to_beta
```

CVS is a very convenient tool for coordinating the work of a large number of developers on a complex project. It provides much of the same functionality as CCSLAND. However, its major strength and its major weakness is that it does not use exclusive locks. It is really a personal choice whether you would prefer to use CCSLAND with its exclusive locks or use CVS with its *copy/modify/merge* strategy. I would choose CCSLAND, but the experts are not unanimous about which is better.

CVS II was written by Brian Berliner and Jeff Polk. It is available through the Free Software Foundation. The distribution tape contains a number of documents which describe CVS functionality in much more detail than what was given here. It should not be confused with its predecessor CVS I which was written by Dick Grune and was released through the *comp.sources.unix* newsgroup.

## Distributed RCS: an even broader solution

Have you ever worked on a project where not all of the project's developers were based at the same site ? Such a working arrangement is becoming increasingly common because of the fact that modern communication and networking technology makes such an arrangement possible. I refer to such a group as a *distributed development team*.

Plain RCS is not capable of managing the sources for a distributed development team because you may have developers at different sites simultaneously working on the same source files. In order to manage the sources for a distributed development team you need a distributed RCS.

Distributed RCS (DRCS), is a system which I developed as part of my Ph.D. project. DRCS allows users at different geographic locations to have the illusion of sharing access to a single set of RCS files. What in fact happens is that there is a replica of each RCS/DRCS file at each of the sites. Whenever a user updates their local replica, the DRCS software transparently propagates this update to all other replicas. Hence, all of the replicas are kept in synchronization with each other, and the users get the illusion of accessing one global RCS file.

The DRCS software consists of, the original RCS code plus some additional routines which I wrote to handle the distributed features of DRCS. The changes which were made to existing RCS routines were minimal; all that was required was that they recognize when to call the new routines. Because DRCS uses the original RCS routines, DRCS is guaranteed to be compatible with RCS. All of the standard RCS commands and options also work with DRCS. Therefore anyone who is currently using RCS can change over to use DRCS without any training. In addition any program that interfaces with RCS, (e.g. CVS or CCSLAND), should not require any modification to interface with DRCS.

The DRCS replica files are almost identical in format to RCS files. The only difference is that, DRCS files contain additional fields to store information about the location and status of the various replicas of the file. Hence, it is possible for DRCS software to read and understand existing RCS files (they are regarded as DRCS files with only one replica). Unfortunately it is not always possible for the standard RCS software to read and understand DRCS files.

By using a distributed file system, (such as NFS), you can easily provide share access to your RCS files from a number of nodes in a local area network. Therefore, it is likely that DRCS will really only be useful for wide area networks where it might not be economically feasible to install communication links capable of supporting a distribute files system. In fact a major strength of DRCS is the fact that it does not require on-line links between each of the sites sharing files. Updates are propagated in the background using any queue oriented communications utility (such as UUCP or mail).

In any distributed system, there is a potential for problems to arise when two users simultaneously make conflicting updates at two different sites. In order to avoid this type of problem, some policy must be adopted for synchronizing updates at different sites. DRCS adopts a simple master/slave type of synchronization policy. The first replica created is called the master replica, all subsequently created replicas are called slave replicas. All updates are first carried out on the master replica and then propagated to the slave replicas. If a user issues an update command on a slave replica, the DRCS software transparently communicates this command to the master site.

While DRCS implements all of the standards RCS commands options, it must also implement a small number of new command options relating to the distributed nature of the system. The most significant new command options are for creating and deleting replicas. New replicas can be created with the **rcs -p** options. Existing replicas can be deleted with the **rcs -P** option.

For example, the following command creates a new replica of the local file *sample.c* at the remote site *london*, the replica file will be located at */usr/src/RCS/sample.c,v* at the *london* site.

```
% rcs -plondon:/usr/src/RCS/sample.c,v
```

Likewise, the following command will delete the existing replica at the *tokyo* site.

```
% rcs -Ptokyo
```

A distributed system such as DRCS can make substantial performance improvements based upon its decision about when/if to propagate updates from the master to the slave sites. The policy that a system uses when deciding when/if to propagate updates is called its replication policy.

Some versions are never read at some of the slave sites, hence it is wasteful to copy over these versions to those slave sites. A replication policy, which decides that versions should only be copied over in response to specific user demand, is called a demand replication policy. A demand replication is very efficient in saving communication costs because versions are never copied over unnecessarily. However, if a demand replication policy is being used, user response is frequently delayed by the need to copy over data from the master site.

The other extreme from demand replication is pre-replication. If a pre-replication policy is being used, all versions are copied to all sites immediately after being created. This replication policy avoids all delays in user response because, whichever version is requested by a user is bound to be stored locally. However, pre-replication is wasteful of communications bandwidth because many versions are copied over to sites that never need them.

RCS/DRCS file contain two different types of data: version attributes and version contents. Version attributes (*e.g.*, the author, creation date etc.) do not comprise much data, hence DRCS uses pre-replication for the version attributes. However, the version contents can comprise a large volume of data, hence DRCS only replicates version contents on demand.

The current DRCS replication policy is very efficient in preserving communications bandwidth (which is expensive in wide area networks). However, the savings in communications bandwidth are often achieved in return for poor performance. Some users might prefer to use a pre-replication policy, because they are willing to pay the increased communications cost in return for improved performance. As a result, a future version of DRCS (which has been developed but not yet released), will allow users to select whichever replication policy they prefer.

DRCS v1.0 was built upon release 4 of RCS, which was acquired under license from Purdue. If you are interested in getting a copy of DRCS v1.0, you can get it by contacting me at the address below. You will be required to sign a form agreeing to comply with the original RCS license terms from Purdue.

Brian O'Donovan,  
c/o Digital Equipment Intl. B.V.,  
Mervue Industrial Estate,  
Galway,  
Republic of Ireland.

Phone : intl - 353 - 91 - 51271  
Fax : intl - 353 - 91 - 53946  
Email : odonovan@ilo.dec.com

However, a more recent version of the RCS code is now available through the Free Software Foundation (GNU). The current release of RCS (v5.6), is available from GNU without the need to sign the Purdue License agreement. As well as being more freely available, this version also has many new features that were not in the version of RCS which was used to build DRCS v1.0. It is very likely that, in the near future a new version of DRCS will be implemented on top of the latest version of RCS. This new version will combine all of the features of DRCS v1.0, RCS v5.6 and some new features which are currently being designed. This version will be distributed through the normal GNU channels.

If you are interested in using DRCS, I would strongly recommend that you should wait until the GNU version becomes available. If you send me a mail message, I will put you on a mail distribution list to receive up to date information about DRCS availability.