

# Experiences in using Optimisitic Locking in Fujaba

## Position paper

Christian Schneider  
SE, Universität Kassel  
Wilhelmshöher Allee 73  
34121 Kassel

christian.schneider@uni-kassel.de

Albert Zündorf  
SE, Universität Kassel  
Wilhelmshöher Allee 73  
34121 Kassel

albert.zuendorf@uni-kassel.de

## 1. REQUIREMENTS ON VERSIONING

In software engineering, we have profound experiences in using versioning systems for textual programming languages. At the beginning we used versioning systems with pessimistic locking like SCCS or RCS. Typically, all source files needed to be read only. When a user wanted to change one file, he or she obtained a write lock (on a local copy). Then, the file was changed, compiled and tested and the change was committed. On larger teams, this pessimistic locking approach runs into problems, when more frequently multiple users need write locks on multiple files. Frequently, some user will have to wait for some other user to release the write lock on a commonly used file.

To overcome these limitations, versioning systems with optimistic locking, such as CVS [5] or SVN [7], have been introduced. These systems allow for concurrent changes on multiple local copies of the same file. The assumption is, that different users will most likely modify disjunct parts of the file and that merging these changes will work smoothly. Before a change can be committed to the central versioning repository, the user has to merge changes already committed by other team members into his local copy. If the different users have modified distinct parts of the file, only, these changes will be combined in the optimistic assumptions that they do not interfere.

If one uses such an optimistic locking system the first time, he or she is usually scared by the idea of merging chunks of changed code into his or her files. However, experiences show that merging works surprisingly smoothly. We use textual versioning with optimistic locking for years now in the Fujaba [6] project with about 30 developers. Merging happens quite frequently. But merge conflicts, i.e. two users have actually changed the same lines of code, are quite seldom. If a merge conflict occurs, some effort is necessary to overcome it. However, the user has just edited these lines of code, so he or she is usually still quite familiar with it and thus merging is bearable.

In case of a merge without a conflict, one may even feel more uncomfortable. However, bugs introduced by non-conflicting merges are quite seldom. The reason for this might be, that team collaboration is organized in a way that changes by different team members do not interfere. In addition, modern programming language have rich compile time checks. In case of interfering changes, the compiler

will most likely detect inconsistencies. However, even this rarely happens.

### 1.1 Vital features

Using optimistic locking enables us to sandbox the team members and their different tasks very effectively. Each team member is able to do his work without being disturbed by changes introduced by other team members until he or she has finished his or her task. After that, merging is usually easy. If problems occur, they are usually caused by some student who has been off duty for several weeks and did not check in major changes before he or she left (or the professor has been on some conferences . . .). After all, from our position, versioning support with optimistic locking is a must have for software development teams.

Due to our experiences, versioning with optimistic locking is indispensable for the development of large UML models, too. Large UML models will most likely be developed by teams of modelers, too. Most likely it is not possible to structure the UML model such that each team member works on a distinct part of the model without the need to modify another's part and without being interfered by modifications on other parts. As far as our experience goes, almost any use case requires changes on multiple classes and methods in multiple packages or components crosscutting multiple architectural layers. Thus, working on a shared model e.g. managed in a common database where each change becomes visible to all team members, instantly, results in lots of disruptions of one team member's work by the work of the others. It will be hard to find a moment in time where the model is consistent e.g. in order to generate code from it.

Locking parts of the model will frequently result in team members blocked by waiting for some other team members to release the lock on some commonly used model parts.

Thus, we strictly vote for versioning support for UML models that provides optimistic locking. Within our UML case tool Fujaba we have developed such an optimistic versioning support, using the CoObRA library (*Concurrent Object Replication framework*, [2, 1, 3]). We use CoObRA for several years now within the development of the Fujaba Tool Suite itself and within student projects. Before we report on our experiences with optimistic locking concepts and merging of UML model changes in the last section, the next section will outline some technical details of the CoObRA

mechanism.

## 1.2 Versioning framework

Fujaba provides a sophisticated versioning support via CoObRA. Basically, CoObRA subscribes itself as listener to all objects of a given model. Then, CoObRA protocols all model changes together with undo information. Additional CoObRA operations allow to group incremental changes from i.e. user actions. Based on these information, CoObRA is already able to provide undo/redo, persistency and versioning functionality.

The CoObRA system is also able to send chunks of change protocols to a project repository. This allows for team support. One user may check in his changes and other users may download these changes into their local version and just replay the change protocol. In order to identify the states of the different local models, CoObRA adds version marks to the change protocols. These version marks allow to identify the changes which have already been downloaded by some user and which have not. Note, CoObRA's versioning is based on object identifiers that are unique over sites and sessions. To achieve this, the server provides unique session ids which are used as prefixes for object ids.

As mentioned, CoObRA provides an optimistic locking concept. On update, the local model is rolled back to the state of the last synchronization. Then, the changes received from the server are applied. This creates no conflicts since the local changes have been rolled back. Next, the local changes are replayed. This might create merge conflicts, e.g. if a locally modified object has been deleted on the server site or if a locally modified attribute has been modified on the server, too. In case of a conflict, all changes of the corresponding user interaction are rolled back and the merge mechanism continues with the next local user interaction. This might result in additional conflicts, e.g. if subsequent changes rely on objects that would have been created by a rolled back user interaction. CoObRA collects all conflicting changes and provides a list of these for the application. The application may then display a user interface which shows the conflicting changes in a user friendly visualization, and which enables the user to resolve the conflict, manually.

The Fujaba Tool Suite uses the CoObRA versioning mechanism in two ways. First, Fujabas code generation provides dedicated support in order to add CoObRA's versioning support to applications modeled with Fujaba. Second, Fujaba uses CoObRA's versioning mechanism for its internal meta model, too. In the latter case, we use a very simple user interface for check in, update, and for the visualization of conflicting changes. Actually, we just use the toString method of changes in order to visualize the list of conflicting changes. Although this visualization refers to internal names of our metamodel, as far as our experience goes, this simple visualization works already sufficiently well for most cases. The next section will report more details on our experiences with the versioning of Fujabas model using CoObRA.

## 2. EXPERIENCES

We use CoObRA's versioning for Fujabas internal model since several years now. We have two main projects with several 10,000 lines of code generated from the models, the

CodeGen2 and the OBA project. The first specifies and realizes and bootstraps Fujabas code generation. The second is an industrial project developing an expert system for the design of wiring harnesses of modern cars. In both projects, the number of developers is relatively small and thus, we observed only very few merge conflicts. In addition, the developers within these projects are very familiar with Fujabas metamodel and thus, it was easy for them to interpret the simple conflict reports provided by CoObRA's standard mechanism.

In addition, we have used CoObRA to version the Fujaba projects of our students in several UML and modeling courses at University of Kassel. This course is run in the 4th or 5th term and it is mandatory for all computer science students. We have about 80 students per course and we organize them in small teams of about 5 students. We run this course as a software development project. Each team has to develop a UML model of some board game with a reasonable complexity. The model has to cover class diagrams and behavior specifications for method bodies and for unit tests. From these specifications, all code is generated, automatically, and the unit tests are run to validate the model behavior.

These student projects are use case driven. This means, each team member considers one use case after the other. For each use case the responsible team mate models a scenario specification, which is later on turned into a unit test. In addition, the responsible team mate extends and adapts the common class diagram as necessary. Finally, the responsible team mate models the functionality that realizes the use case, generates code, and validates his model. When a use case has been accomplished, the responsible team mate synchronizes the content of his local working copy with the team repository using CoObRA.

Within these student projects, the team members most likely interfere within the common class diagram. Accordingly, we observe a certain number of merge conflicts due to renaming and deletion of common items. Another kind of conflict occurs, if two team members introduce a new class or attribute or association for the same purpose, independently. For CoObRA, this is no conflict and the merge works smoothly. Actually, the team members have to unify these model elements, manually, in order to achieve a collaboration of their individual contributions.

### 2.1 Findings

After all, we are very happy with CoObRA's optimistic versioning support. Team members are sandboxed as we are used to it with textual programming languages. Most of the time, merging runs smoothly. This holds especially for larger models where the chances for a concurrent modification of the same model item decreases. Maybe this is also facilitated due to our use case driven modeling approach where working packages overlap mainly in the common class diagram which becomes relatively stable after some project development time.

However, merging causes a number of problems, also. First of all, if a merge conflict occurs, Fujaba provides only very cryptic reports on refused changes, currently. Astonishingly, even our students are able to deal with our cryptic reports

quite well. For example, they interpret the message

```
alter field de.fujaba.uml.UMLClass("Student").name  
from "Person" to "Collegian" failed
```

easily, since they usually recall that they just have changed the name of the “Person” class to the new name “Collegian” and some team mate has probably renamed the same class to “Student”.

Of course, merge conflicts may result in model inconsistencies. For example, one may have added a lot of object instances of some Class C to his or her model and some team mate may have removed class C. In Fujaba this results in a merge conflict since the instanceOf link connecting the object instance and its declaring class cannot be inserted into the merged project. Therefore, all these object insertions are refused. Luckily, the deletion of class diagram elements get more seldom when the project matures.

Merging may also generate context sensitive model inconsistencies. For example, two team mates may concurrently introduce a new class with the same name e.g. Person. Since CoObRAs versioning is based on internal ids, after merging the project will contain two classes with the same name. In general, merging may result in model states that could not be reached by usual editing. Any consistency constraint that is usually maintained by the user interface may be violated. Thus, merging sometimes corrupts the internal object structures of our tool. To deal with these kinds of problems, we steadily enhance our metamodel and tool in order to allow temporary inconsistencies. This would e.g. allow temporary name conflicts or missing declarations. This is desirable even without any versioning to foster intuitive editing for the user. As the consistency of the model still needs to be ensured, measures to fix violated consistency constraints should be provided lateron. Currently, we provide some repair mechanisms that resolve typical model inconsistencies on request. Finally, our dynamic object browser eDOBS enables our experts to repair any internal model problems.

### 3. CONCLUSIONS

After all, we strongly believe that versioning of UML models is absolutely necessary. This versioning must provide optimistic locking mechanisms and merging mechanisms as it is state-of-the-art for textual documents. Our experiences with Fujaba and CoObRA show that versioning of UML and other object oriented models with optimistic locking is feasible. Of course, there are a number of pitfalls. For using CoObRA the metamodel must provide means to protocol changes, e.g. a notification mechanism. The metamodel must not rely on too many additional consistency constraints, or must be able to deal with otherwise refused changes. Finally, merging may result in conflicts, some kinds of inconsistencies, and editing operations may get lost due to such problems and the project may reach inconsistent states that could not be reached via usual editing operations. However, due to our experiences the benefits of being able to have local working copies for the members of a large project team, working on a common UML modeling project, outweigh the drawbacks several times. Within Fujaba, we can just not imagine, how someone can live without versioning support and optimistic locking for his or her models.

### 4. REFERENCES

- [1] C. Schneider. CASE Tool Unterstützung für die Delta-basierte Replikation und Versionierung komplexer Objektstrukturen (Diploma Thesis, german), 2003.
- [2] C. Schneider, A. Zündorf, and J. Niere. CoObRA - a small step for development tools to collaborative environments. In *Workshop on Directions in Software Engineering Environments; 26th international conference on software engineering*. ICSE 2004, Scotland, 2004.
- [3] CoObRA 2. <http://www.se.eecs.uni-kassel.de/se/?coobra>, 2006.
- [4] CoObRA Server Control. <http://www.se.eecs.uni-kassel.de/se/?coobras>, 2006.
- [5] CVS - Concurrent Versions System. <http://www.nongnu.org/cvs/>, 2006.
- [6] The Fujaba Toolsuite. <http://www.fujaba.de/>, 2006.
- [7] SVN - Subversion. <http://subversion.tigris.org/>, 2006.